

A complex, abstract graphic composed of numerous small, glowing blue and purple dots connected by thin lines, forming a network or cloud-like structure against a dark background. In the center, the words 'Build.', 'Unify.', and 'Scale.' are overlaid in large, bold, white and green letters.

Build. Unify. Scale.

WIFI SSID:SparkAISSummit | Password: UnifiedAnalytics

ORGANIZED BY
 databricks



SPARK+AI
SUMMIT 2019

Spark Streaming

Headaches and Breakthroughs in Building
Continuous Applications

Landon Robinson & Jack Chapa

SPOTX

#UnifiedAnalytics #SparkAISummit

Who We Are



Landon Robinson
Data Engineer



Jack Chapa
Data Engineer

Big Data Team @ SpotX

But first... why are we here?

Because Spark Streaming...

- is very **powerful**
- can **supercharge your infrastructure**
- ... and *can* be very complex!

Lots of headaches and breakthroughs!

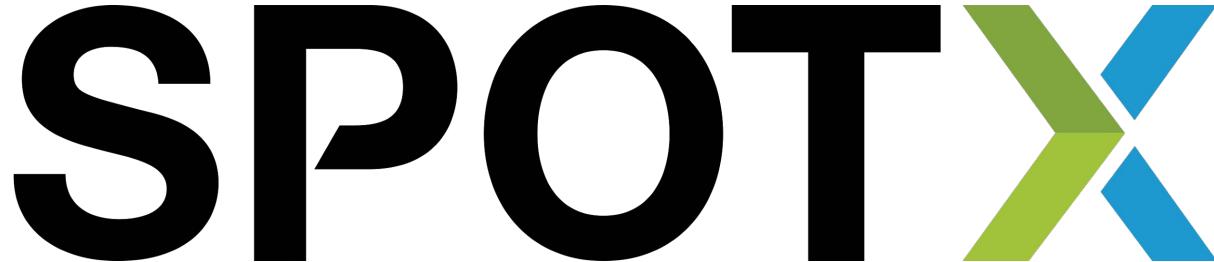
Takeaways

Leave with a few actionable items that we wish we knew when we started with Spark Streaming.

Focus Areas

- Streaming Basics
- Testing
- Monitoring & Alerts
- Batch Intervals & Resources
- Helpful Configurations
- Backpressure
- Data Enrichment & Transformations

Our Company



The Trusted Platform For
Premium Publishers and
Broadcasters

We Process a Lot of Data



Data:

- 220 MM+ Total Files/Blocks
- 8 PB+ HDFS Space
- 20 TB+ new data daily
- 100MM+ records/minute
- 300+ Data Nodes

Apps:

- Thousands of daily Spark apps
- Hundreds of daily user queries
- Multiple 24/7 Streaming apps

Spark Streaming is Key for Us



Our uses include:

- **Rapid ingestion of data** into warehouse for querying
- **Machine learning on near-live data streams**
- **Ability to react to and impact live situations**
- **Accelerated processing / updating** of metadata
- **Real-time visualization** of data streams and processing



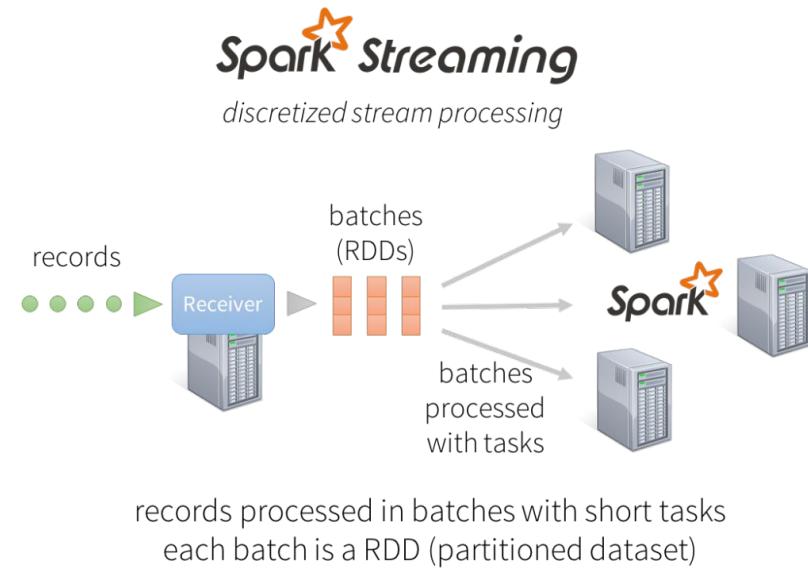
Spark Streaming Basics

a brief overview

Spark Streaming Basics

Spark Streaming is an extension of Spark that enables scalable, high-throughput, fault-tolerant processing of live data streams.

- **Stream ==** live data stream
 - **Topic ==** Kafka's name for a stream
-
- **DStream ==** sequence of RDDs formed from reading a data stream
- **Batch ==** a self-contained job within your Streaming app that processes a segment of the stream.





Testing

*Rapid development and
testing of Spark apps*

Use Spark in Local Mode

You can start building Spark Streaming apps in minutes, **using Spark locally!**

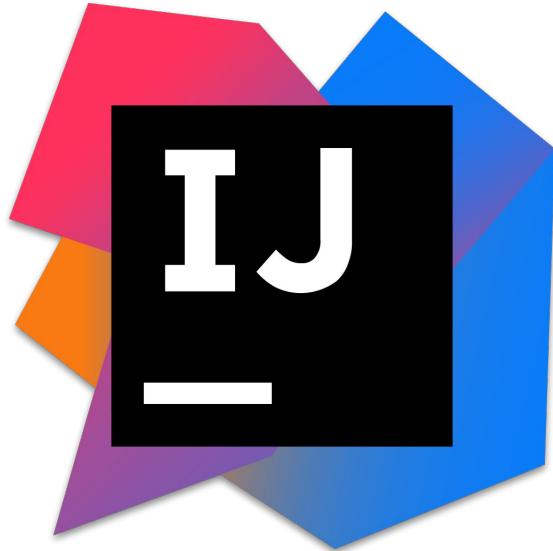
On your local machine

- No cluster needed!
- Great for rough testing

We Recommend:

IntelliJ Community Edition

- with SBT: For dependency management



Use Spark in Local Mode

The Scala Build Tool is your friend!

Simply:

- Import Spark libraries
- Invoke a Context and/or Session
- Set master to local[*] or local[n]

In your build.sbt:

- src/test/scala => “provided”
- src/main/scala => “compiled”

```
// Dependency Chain
libraryDependencies ++= Seq(
  "org.apache.spark" %% "spark-core" % spark_version % "provided",
  "org.apache.spark" %% "spark-sql"   % spark_version % "provided",
  "org.apache.spark" %% "spark-hive"  % spark_version % "provided",
```

Example Unit Test using just a SparkContext

```
// Setup Spark
val spark = SparkSession
  .builder()
  .master("local[1]")
  .appName("CassandraTests")
  .enableHiveSupport()
  .getOrCreate()

// Import Implicits (case class casting)
import spark.implicits._

// Unit Test #1: Retrieve audience_data sample and cast it
"Cassandra Connector" should "retrieve a 10 row audience_data sample from Cassandra (dev env)" in {
    val df = spark
      .read
      .format("org.apache.spark.sql.cassandra")
      .options(table_audience)
      .load()
      .as[Audience]
      .limit(10)

    assert(df.count() === 10L)
}
```

Invoke a local session:

- In your unit test classes
- Test logic on small datasets

**Add to your deployment pipeline
for a nice pre-release gut check!**

Unit Testing

Spark Streaming Apps can easily be unit tested

- Using `.queueStream()`
- Using a spark testing library

Libraries

- spark-testing-base
- sscheck
- spark-tests

Use Cases

- DStream actions
- Business Logic
- Integration

Example Library: spark-testing-base

- Easy to Use
- Helpful wrappers
- Integrates w/ *scalatest*
- Minimal code required
- Clock management
- Runs alongside other tests

```
class FileSparkStreamingTests extends FunSuite
  with StreamingSuiteBase with SharedSparkContext {

  test(testName = "Convert CSV data to Session") {
    val testData: List[List[String]] = List(FileSparkStreamingTests.getTestData)
    val expected = List(List(FileSparkStreamingTests.getSession))

    def validateResults(ds: DStream[String]): DStream[Session] = {
      FileSparkStreaming.processSessionData(ds)
    }

    testOperation[String, Session](testData, validateResults _, expected, ordered = true)
  }
}

val spark_version = "2.4.0"
val spark_testing_version = spark_version + ".0.11.0"

libraryDependencies += Seq(
  "org.apache.spark" %% "spark-core" % spark_version % "provided",
  "org.apache.spark" %% "spark-sql" % spark_version % "provided",
  "org.apache.spark" %% "spark-streaming" % spark_version % "provided",
  "com.typesafe.slick" %% "slick" % "3.3.0",
  "org.slf4j" % "slf4j-nop" % "1.7.10",
  "com.holdenkara" %% "spark-testing-base" % spark_testing_version % "test"
)

fork in Test := true
javaOptions += Seq("-Xms512M", "-Xmx2048M", "-XX:MaxPermSize=2048M", "-XX:+CMSClassUnloadingEnabled")
```



GitHub: <https://github.com/holdenk/spark-testing-base>



Monitoring

*Tracking and visualizing
performance of your
app*

Monitoring is Awesome

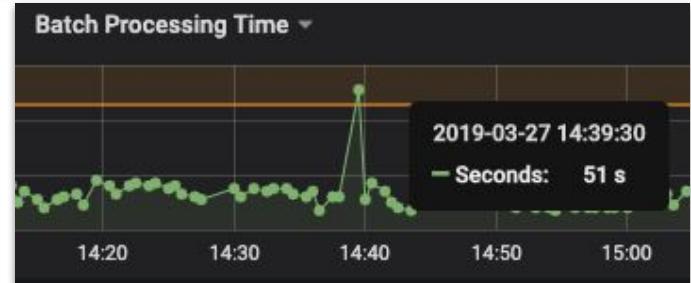
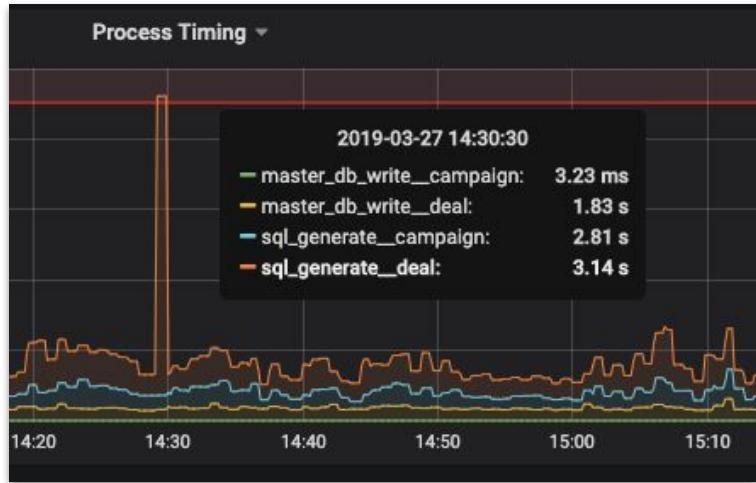
It can reveal:

- How your app is performing
- Problems + Bugs!

And provide opportunities to:

- See and address issues
- Observe behavior visually

But monitoring can be *tough* to implement!



Monitoring (a less than ideal approach)

You *could* do it all in the app...

Example: Looping over RDDs to:

- Count records
- Track Kafka offsets
- Processing time / delays

But it's less than ideal...

- Calculating performance significantly impacts performance... not great.
- All these metrics are **calculated by Spark!**

```
myStreamingData.count().print()
```

```
myStreamingData.foreachRDD(rdd => {  
    rdd.map(x => {  
        println(x.topic())  
        println(x.offset())  
        println(x.partition())  
    })  
})
```

```
val t0 = System.nanoTime()  
// execute a task  
val t1 = System.nanoTime()  
val processing_time = (t1 - t0) / 1000  
println("I took this long: " + processing_time)
```

Monitoring and Visualization (using Listeners)

Use **Spark Listeners** to access metrics in the background!

Let Spark do the hard work:

- Batch duration, delays
- Record throughput
- Stream position recovery

**Come to our talk: Spark Listeners:
A Crash Course in Fast, Easy
Monitoring!**

- Room 3016 | Today @ 5:30 PM

```
package org.apache.spark.streaming.scheduler
+
@DeveloperApi
trait StreamingListener {

    /** Called when a receiver has been started */
    def onReceiverStarted(receiverStarted: StreamingListenerReceiverStarted) { }

    /** Called when a receiver has reported an error */
    def onReceiverError(receiverError: StreamingListenerReceiverError) { }

    /** Called when a receiver has been stopped */
    def onReceiverStopped(receiverStopped: StreamingListenerReceiverStopped) { }

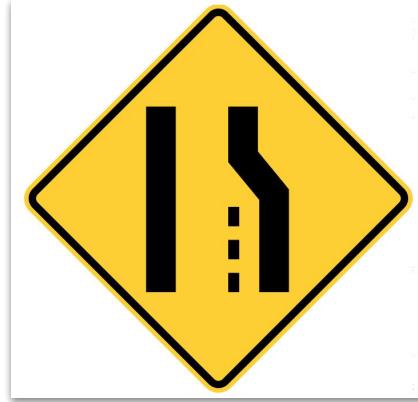
    /** Called when a batch of jobs has been submitted for processing. */
    def onBatchSubmitted(batchSubmitted: StreamingListenerBatchSubmitted) { }

    /** Called when processing of a batch of jobs has started. */
    def onBatchStarted(batchStarted: StreamingListenerBatchStarted) { }

    /** Called when processing of a batch of jobs has completed. */
    def onBatchCompleted(batchCompleted: StreamingListenerBatchCompleted) { }

    /** Called when processing of a job of a batch has started. */
    def onOutputOperationStarted(
        outputOperationStarted: StreamingListenerOutputOperationStarted) { }

    /** Called when processing of a job of a batch has completed. */
    def onOutputOperationCompleted(
```



Kafka Offset Recovery

*Saving your place
elsewhere*

Writing Offsets to MySQL

Inside the Spark
Listener class, after a
batch completes...

You can access an
object generated by
Spark containing your
offsets processed.

Take those offsets and
back them up to a
DB...

```
/**  
 * Takes a topic object and writes the max offset for each partition it contains this batch to MySQL.  
 *  
 * @param topic A topic object within a Batch's StreamIdToInputInfo  
 */  
def writeTopicOffsetsToMySQL(topic: Tuple2[Int, StreamInputInfo]): Unit = {  
  
    // map offset info to OffsetRange objects  
    val partitionOffsets = topic._2.metadata("offsets").asInstanceOf[List[OffsetRange]]  
  
    // for every partition's range of offsets  
    partitionOffsets.map(offsetRange => {  
  
        // write the new starting offset for each partition in the topic to the state db  
        var maxOffset = offsetRange.untilOffset - 1  
  
        // create a now() timestamp  
        val now = new DateTime().toString("YYYY-MM-dd HH:mm:ss")  
  
        // form the sql  
        val sql =  
            s"""INSERT INTO $mySQLDB.$mySQLTable (consumer, topic, partition_id, offset, offset_ts, batch_size)  
                VALUES  
                ('$mySQLConsumer', "${offsetRange.topic}", ${offsetRange.partition}, '$maxOffset', '$now', ${offsetRange.count})  
                ON DUPLICATE KEY UPDATE offset_ts = VALUES(offset_ts), offset = VALUES(offset),  
                batch_size = VALUES(batch_size)  
                """  
  
        // execute the sql to offload offsets to the table  
        val st = mySQLConnectionPool.createStatement  
        st.execute(sql)  
        st.close()  
    })  
}
```

Reading Offsets from MySQL

Your offsets are now stored in a DB after each batch completes.

Whenever your app restarts, it reads those offsets from the DB...

And starts processing where it last left off!

```
def subscribeToTopicWithState(  
    streamingContext: StreamingContext,  
    topic: String,  
    kafkaParams: Map[String, Object],  
    dbConnection: MySQLConnection,  
    stateConsumer: String,  
    resetOffsets: Boolean = false  
): InputDStream[ConsumerRecord[String, Option[GenericRecord]]] = {  
  
    // Store list of topics in an array  
    val topics = Array(topic)  
  
    // Either Get Offsets from a State Database or Reset to Latest  
    val offsets = getOffsetsFromState(dbConnection, stateConsumer, topic, resetOffsets)  
  
    // Subscribe to the topic based on the offsets provided  
    val strategy = getStrategy(offsets, kafkaParams, topics)  
  
    // Create a direct stream and pull data  
    val subscribed = KafkaUtils.createDirectStream[String, Option[GenericRecord]](  
        streamingContext,  
        LocationStrategies.PreferConsistent,  
        strategy  
    )  
  
    // Filter data coming out of Kafka to remove bad records  
    subscribed  
}
```

Getting Offsets from the Database

```
def getOffsetsFromState(dbConnection: MySQLConnection,
                      consumer: String,
                      topic: String,
                      resetOffsets: Boolean
                     ): Map[TopicPartition, Long] = {

    // Establish a MySQL Connection given a URL
    val stateDbConn = ConnectionPool(dbConnection.toString).getConnection

    // If You want to reset offsets, just return an empty map
    val offsets =
        if (resetOffsets) {
            new HashMap[TopicPartition, Long]()
        }

    // If you want to use your existing offsets, get them from MySQL
    else {
        val stmt = stateDbConn.createStatement
        val query = "SELECT * FROM " + dbConnection.database + "." + dbConnection.table +
                    " WHERE consumer = '" + consumer + "' AND topic = '" + topic + "'"
        logger.info(query)
        val result = stmt.executeQuery(query)
        new RsIterator(result).map(
            x =>
                new TopicPartition(
                    x.getString("topic"),
                    x.getInt("partition_id")
                ) ->
                    (x.getLong("offset") + 1L) // Add one because we record last known offset
            ).toMap
    }

    offsets
}
```

Example: Reading Offsets from MySQL

```
def subscribeToTopicWithState(
    streamingContext: StreamingContext,
    topic: String,
    kafkaParams: Map[String, Object],
    dbConnection: MySQLConnection,
    stateConsumer: String,
    resetOffsets: Boolean = false
): InputDStream[ConsumerRecord[String, Option[GenericRecord]]] = {

    // Store list of topics in an array
    val topics = Array(topic)

    // Either Get Offsets from a State Database or Reset to Latest
    val offsets = getOffsetsFromState(dbConnection, stateConsumer, topic, resetOffsets)

    // Subscribe to the topic based on the offsets provided
    val strategy = getStrategy(offsets, kafkaParams, topics)

    // Create a direct stream and pull data
    val subscribed = KafkaUtils.createDirectStream[String, Option[GenericRecord]](
        streamingContext,
        LocationStrategies.PreferConsistent,
        strategy
    )

    // Filter data coming out of Kafka to remove bad records
    subscribed
}
```

Example: Reading Offsets from MySQL

```
def getStrategy(
    offsets: Map[TopicPartition, Long],
    kafkaParams: Map[String, Object],
    topics: Array[String]
): ConsumerStrategy[String, Option[GenericRecord]] ={

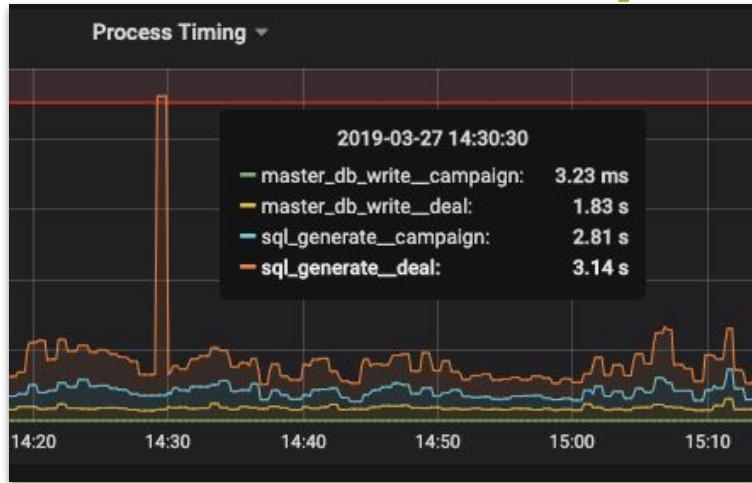
    // if offsets were pulled from MySQL, assign the stream to start there
    val strategy = if (offsets.size > 0) {
        ConsumerStrategies.Assign[String, Option[GenericRecord]](offsets.keys.toList, kafkaParams, offsets)
    }

    // otherwise, start at latest data
    else {
        ConsumerStrategies.Subscribe[String, Option[GenericRecord]](topics, kafkaParams)
    }

    strategy
}
```

Timing Logging (around actions)

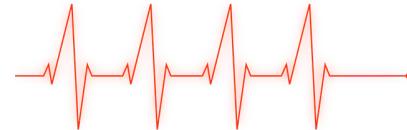
- Record timing info for fast troubleshooting
 - Escalate alarms to the appropriate team
 - Quickly resolve while app continues running



React

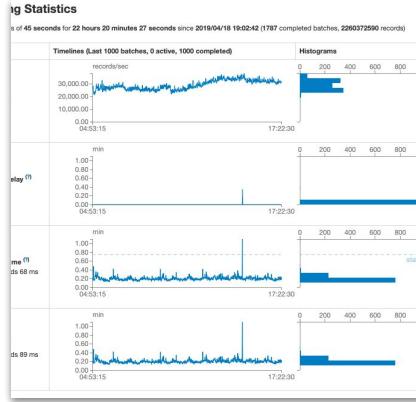
How do I react to this monitoring?

- Heartbeats
- Scheduled Monitor Jobs
 - Version Updates
 - Ensure Running
 - Act on failure/problem
- Monitoring Alarms
- Look at them!



pagerduty





Batch Intervals

*Optimizing for speed
and resource efficiency*

Setting Appropriate Batch Intervals

An **appropriate batch interval** is key to an app that is quick and efficient.

```
def main(args: Array[String]): Unit = {  
  
    val sparkConf = new SparkConf()  
        .setAppName("My Streaming App")  
        .setMaster("local[*]")  
  
    val ssc = new StreamingContext(sparkConf,  
                                batchDuration = Seconds(30))
```



Batch interval

Effectiveness of interval is affected by:

- Resources alloc (cpu + ram)
- Quantity of work
- Quantity of data

You want batches that **process faster than the interval**, but **not so fast that resources are idling and therefore wasted!**

Setting Appropriate Batch Intervals

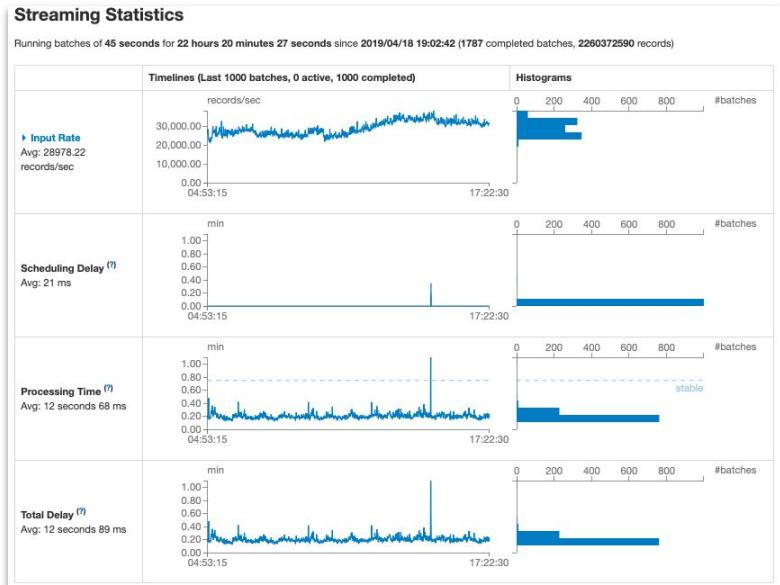
Consider these questions:

How quickly do I need to process data?

- Can I slow it down to save resources?

What is my resource budget / allocation?

- Can I increase? Can I cut back?
- **Bigger interval = more time to process**
- ... but also more data to process
- **Smaller interval = the opposite**



Setting Appropriate Batch Intervals

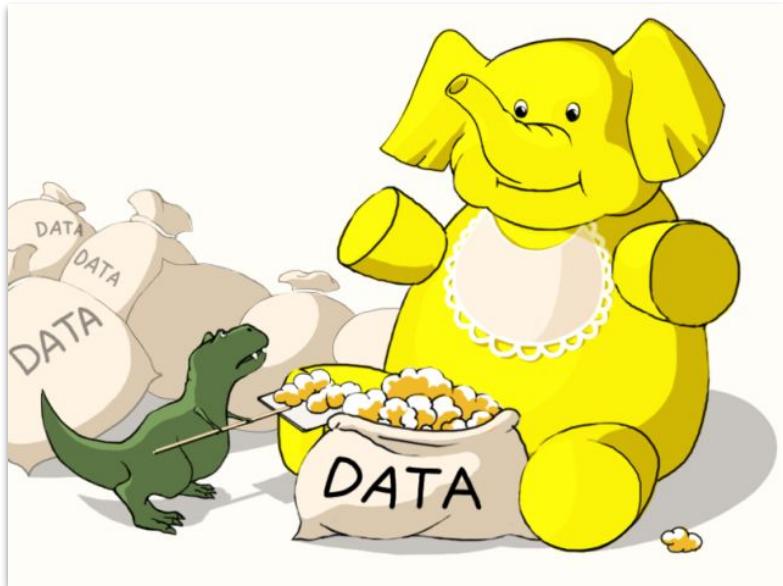
Tips for finding an optimal combination:

Start small!

- a. **Short batch interval (seconds)**
- b. **Modest resources**

Whichever you have in more flexible supply (a or b), increasing accordingly.

Again: processing time < interval = good
Comfortably less, not significantly less.



Additional Resource Notes

- **Scale down** when possible
 - Free up resources or save on cloud utilization spend
- **Avoid preemption**
 - Use resource pools with prioritization
 - With preemption disabled if you can
- **Set appropriate # of partitions** for Kafka topics
 - Higher volume == higher partition count
 - Higher partition count == greater parallelization



Helpful Configuration Settings

*Configuring your app to
be performant and
efficient*

Helpful Configuration Settings

Spark

- `spark.memory.useLegacyMode = true`
 - `spark.storage.memoryFraction=0.03`
- `spark.submit.deployMode = cluster`
- `spark.serializer = org.apache.spark.serializer.KryoSerializer`
- `spark.rdd.compress = true`
 - `spark.io.compression.codec=org.apache.spark.io.SnappyCompressionCodec`
- `spark.shuffle.service.enabled = true`
- `spark.streaming.blockInterval = 300`

Kafka

- `enable.auto.commit = 'false'`

Backpressure

Use Case:

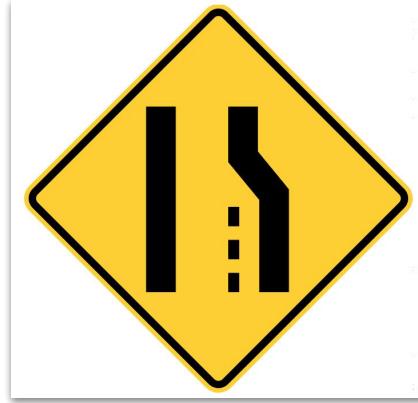
You have irregular spikes in message throughput from Kafka topics

- Backpressure dynamically alters rate data is received per batch from Kafka.
- Prevents overwhelming of app at startup and peak load.

Settings:

- `spark.streaming.backpressure.enabled = true`
- `spark.streaming.kafka.maxRatePerPartition = 20000`
 - max rate (messages/second) at which each Kafka partition will be read
- **PID Rate Estimator:** can be used to tweak the rate based on batch performance
 - `spark.streaming.backpressure.pid.*`

Source: <https://www.linkedin.com/pulse/enable-back-pressure-make-your-spark-streaming-production-lan-jiang/>



Transformations

*Bringing streaming and
static data together*

Transformations (Streaming + Static)

`transform()`

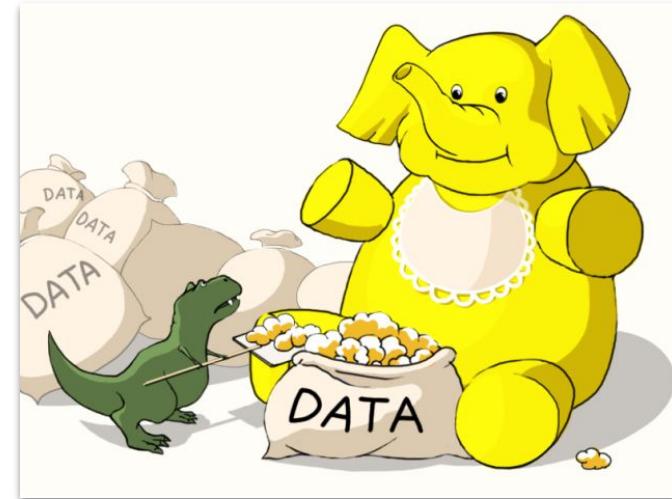
- Allows RDD-level access to data.
- Use case: joining with another RDD

`updateStateByKey() / mapWithState()`

- Apply function to each key - useful for keeping track of state
- Use case: maintaining state between batches (e.g. rolling join w/ two streams)

`reduceByKey()`

- Reduce a keyed RDD with appropriate function.
- Use case: deduping, aggregations



Joining Streaming and Static Data

Using the transform() method on DStream:

Apply an RDD-to-RDD function to every RDD of the DStream. Used for arbitrary RDD operations on the DStream.

- **Useful for applying arbitrary RDD operations on a DStream.**
- **Great for enriching streaming data with supplemental static data**

```
transactions = ... // streaming dataset (dstream)
```

```
transaction_details = ... // static dataset (rdd)
```

```
val complete_transaction_data = transactions.transform(live_transaction =>
  live_transaction.join(transaction_details))
```

Source: <https://hadoopsters.net/2017/11/26/how-to-join-static-data-with-streaming-data-dstream-in-spark/>

Effective Static Joining

How do we handle static and persistent data?

Driver:

- Broadcast if small enough
- Read on driver every batch, then join

Worker:

- Connect on worker - lazy val connection object
- Useful for persisting data

Review

Streaming isn't always easy... but **here are some great takeaways!**

- **Testing:** Use Spark Locally w/ Unit Tests
- **Monitoring:** Use Listeners & React
- **Batch Intervals & Resources:** Be thoughtful!
- **Configuration:** Lots of awesome ones!
- **Transformations:** Do more with your streaming data!
- **Offset Recovery:** Stop worrying and love the offset management!

Contact Us

Landon Robinson

- lrobinson@spotx.tv

Jack Chapa

- jchapa@spotx.tv

hadoopsters.dev

<https://gist.github.com/hadoopsters>



SPARK+AI
SUMMIT 2019

Q & A

#UnifiedAnalytics #SparkAISummit



SPARK+AI
SUMMIT 2019

DON'T FORGET TO RATE
AND REVIEW THE SESSIONS

SEARCH SPARK + AI SUMMIT



Download on the
App Store



GET IT ON
Google Play

