

A dark blue background featuring a complex network of glowing purple and teal dots connected by thin lines, forming a grid-like structure. Interspersed among the dots are binary code sequences (0s and 1s). Overlaid on this graphic is the text 'Unify.' in a large, bold, lime-green font, and 'Build.' and 'Scale.' in a large, bold, white font.

Unify.  
Build.  
Scale.

WIFI SSID:SparkAISSummit | Password: UnifiedAnalytics

ORGANIZED BY  
 databricks



SPARK+AI  
SUMMIT 2019

# Spark Listeners

## A Crash Course in Fast, Easy Monitoring

Landon Robinson & Ben Storrie

**SPOTX**

#UnifiedAnalytics #SparkAISummit

# Who We Are



**Landon Robinson**  
*Data Engineer*



**Ben Storrie**  
*Data Engineer*

**Big Data Infrastructure Team @ SpotX**

# But first... why are we here?

- Building monitoring for large apps can be **tough**.
- You want an **effective and reliable solution**.
- **Spark Listeners are an awesome solution!**

# Our Company



(we show you video ads)  
(which means we also process a lot of data)

# We Process a Lot of Data



## Data:

- 220 MM+ Total Files/Blocks
- 8 PB+ HDFS Space
- 20 TB+ new data daily
- 100MM+ records/minute
- 300+ Data Nodes

## Apps:

- Thousands of daily Spark apps
- Hundreds of daily user queries
- Several 24/7 Streaming apps

# Spark Streaming is Key for Us



Spark Streaming apps are **critical** for us.

They:

- a) **process billions of records** every day
- b) need **real-time monitoring** (*visual dashboards + alerts*)
- c) *without sacrificing performance*

So... a clever solution is necessary.

# Our Solution

Monitoring and visualization of a Spark Streaming app that is:

- Fast
- Easy to integrate
- Has zero-to-minimal performance impact

# Goal & Talking Points

***Equip you with the knowledge and code to get fast, easy monitoring implemented in your app using Spark Listeners.***

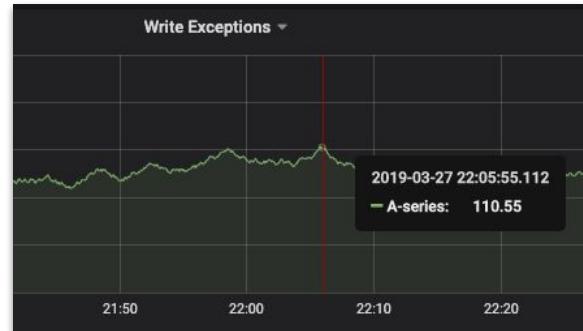
## Talking Points

- Common Spark Streaming Metrics
- Monitoring is Awesome
- Monitoring can be Difficult
- Less than Ideal Monitoring
- Better Monitoring
  - Spark Listeners!
    - StreamingListeners
    - BatchListeners

# Common Metrics of Spark Streaming Apps

## Performance:

- batch duration
- batch scheduling delays
- # of records processed per batch



## Problems:

- failures & exceptions

## Recovery:

- Position (offset) within Kafka topic

```
mysql> select * from offsets where app='myApp1' limit 10;
+-----+-----+-----+-----+-----+-----+-----+
| app | topic | partition_id | offset | offset_ts | batch_size | modified_date |
+-----+-----+-----+-----+-----+-----+-----+
| myApp1 | kafka.some_topic1 | 0 | 877744060 | 2019-03-27 23:30:55 | 836 | 2019-03-27 23:30:55 |
| myApp1 | kafka.some_topic1 | 1 | 877194072 | 2019-03-27 23:30:55 | 836 | 2019-03-27 23:30:55 |
| myApp1 | kafka.some_topic1 | 2 | 877804618 | 2019-03-27 23:30:55 | 810 | 2019-03-27 23:30:55 |
| myApp1 | kafka.some_topic1 | 3 | 877733267 | 2019-03-27 23:30:55 | 785 | 2019-03-27 23:30:55 |
| myApp1 | kafka.some_topic1 | 4 | 878811924 | 2019-03-27 23:30:55 | 859 | 2019-03-27 23:30:55 |
| myApp1 | kafka.some_topic1 | 5 | 877818167 | 2019-03-27 23:30:55 | 813 | 2019-03-27 23:30:55 |
+-----+-----+-----+-----+-----+-----+-----+
```

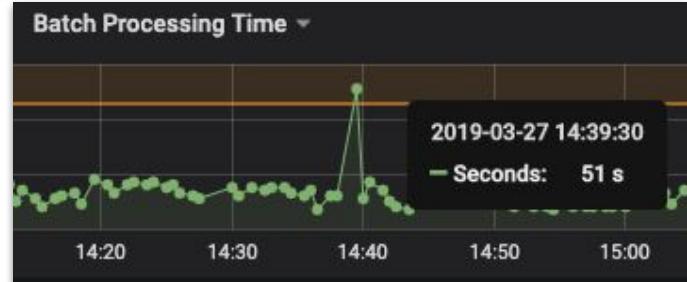
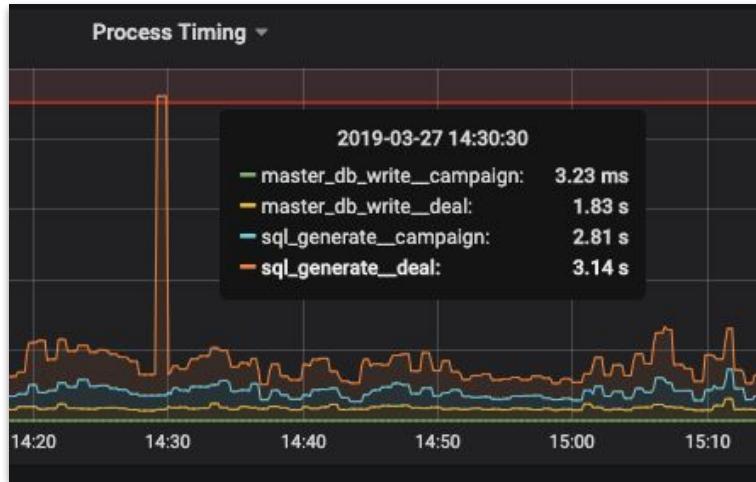
# Monitoring is Awesome

It can reveal:

- How your app is performing
- Problems + Bugs!

And provide opportunities to:

- See and address issues
- Observe behavior visually



# Monitoring can be Difficult to Implement (especially when working with very big data)

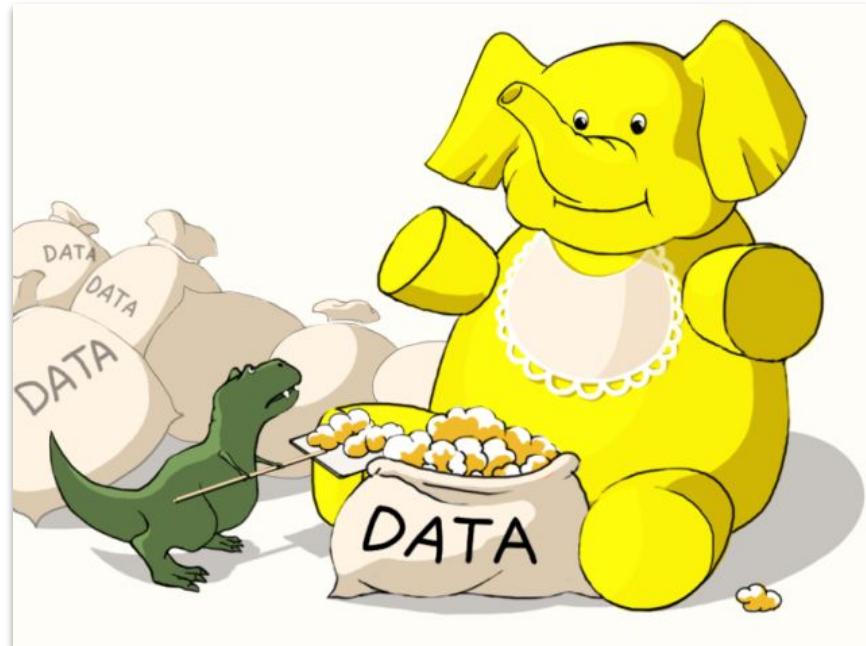
Usually because it requires reading your data!

**Example: `.count().print()`**

- Shuffling involved, and is often a blocking operation!

Mapping over data to gather metrics can:

- be expensive
- and add processing delays
  - Virtually unreliable with big data



# Less than Ideal Monitoring

Make the app compute *all* your metrics!

Example: Looping over RDDs to:

- Count records
- Track Kafka offsets
- Processing time / delays

Why is this bad?

- Calculating performance significantly impacts performance... not great.
- All these metrics are **calculated by Spark!**

```
myStreamingData.count().print()
```

```
myStreamingData.foreachRDD(rdd => {  
    rdd.map(x => {  
        println(x.topic())  
        println(x.offset())  
        println(x.partition())  
    })  
})
```

```
val t0 = System.nanoTime()  
// execute a task  
val t1 = System.nanoTime()  
val processing_time = (t1 - t0) / 1000  
println("I took this long: " + processing_time)
```

**There has to be a better way...**



**Streaming  
Listeners!**

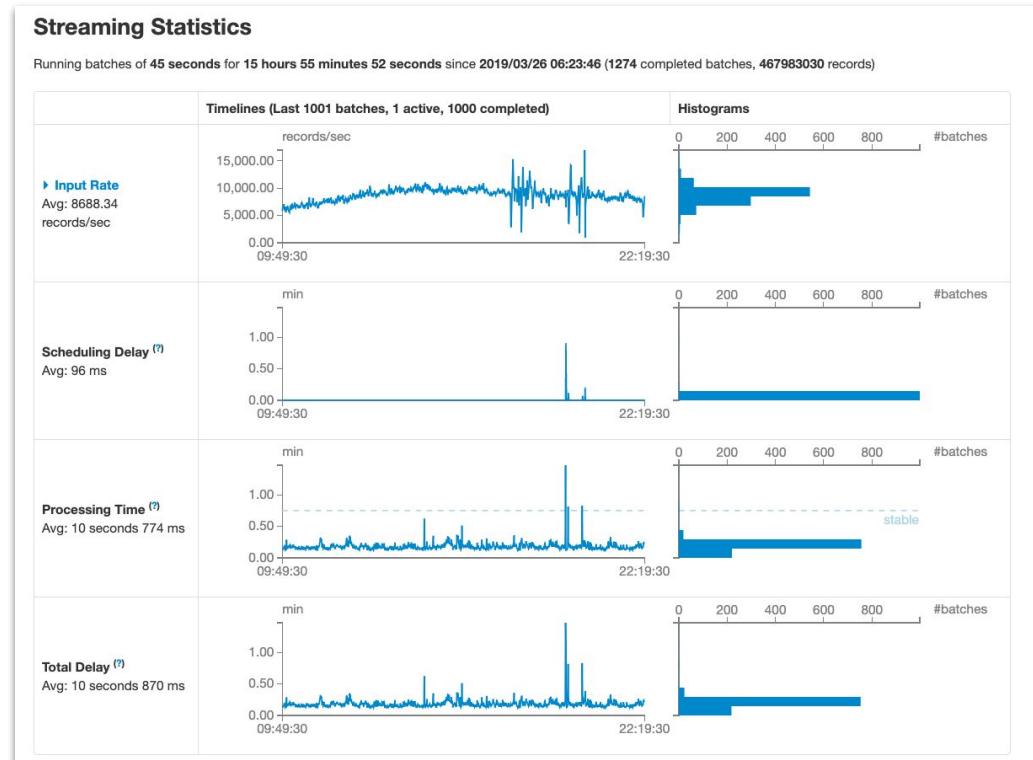
# The Better Way: Streaming Listeners

Let a Spark Streaming Listener handle the basics metrics.

Only make your app compute the “special” metrics it has to.

Spark already crunches some numbers for you -- most evident in the Streaming tab.

So let's take advantage of it!



# Streaming Listeners: *StreamingListener Trait*

The *StreamingListener Trait* within the Developer API has everything you need.

**package** org.apache.spark.streaming.scheduler

It has 8 convenience methods you can override, each with relevant, contextual information from Spark.

**onBatchCompleted()** is the most useful to us. It allows you to execute logic after each batch.

```
/**  
 * :: DeveloperApi ::  
 * A listener interface for receiving information about an ongoing streaming  
 * computation.  
 */  
@DeveloperApi  
trait StreamingListener {  
  
  /** Called when a receiver has been started */  
  def onReceiverStarted(receiverStarted: StreamingListenerReceiverStarted) { }  
  
  /** Called when a receiver has reported an error */  
  def onReceiverError(receiverError: StreamingListenerReceiverError) { }  
  
  /** Called when a receiver has been stopped */  
  def onReceiverStopped(receiverStopped: StreamingListenerReceiverStopped) { }  
  
  /** Called when a batch of jobs has been submitted for processing. */  
  def onBatchSubmitted(batchSubmitted: StreamingListenerBatchSubmitted) { }  
  
  /** Called when processing of a batch of jobs has started. */  
  def onBatchStarted(batchStarted: StreamingListenerBatchStarted) { }  
  
  /** Called when processing of a batch of jobs has completed. */  
  def onBatchCompleted(batchCompleted: StreamingListenerBatchCompleted) { }  
  
  /** Called when processing of a job or a batch has started. */  
  def onOutputOperationStarted(  
    | outputOperationStarted: StreamingListenerOutputOperationStarted) { }  
  
  /** Called when processing of a job of a batch has completed. */  
  def onOutputOperationCompleted(  
    | outputOperationCompleted: StreamingListenerOutputOperationCompleted) { }  
}
```

# Streaming Listeners: Extend Trait in New Listener

Create a new listener that extends the *StreamingListener* Trait.

Override the convenience methods.

Within each is contextual monitoring information.

Let's look at `onBatchCompleted()`.

```
/**  
 * :: SpotXSparkStreamingListener ::  
 * A simple StreamingListener that logs summary statistics across Spark Streaming batches; inherits from DeveloperAPI.  
 *  
 * @param influxHost      Hostname of the Influx service  
 * @param influxDB        Database name in Influx to write to  
 * @param influxMeasurement Measurement name in Influx to write to  
 * @param mySQLHost       Hostname of the MySQL service  
 * @param mySQLDB         Database name in MySQL to write to  
 * @param mySQLTable      Table name in MySQL to write to  
 * @param mySQLUser       Username for authentication in MySQL  
 * @param mySQLPwd        Password for authentication in MySQL  
 * @param mySQLConsumer   Unique name for tracking offsets across streaming apps  
 */  
class SpotXSparkStreamingListener (influxHost: String,  
                                    influxDB: String,  
                                    influxMeasurement: String,  
                                    mySQLHost: String,  
                                    mySQLDB: String,  
                                    mySQLTable: String,  
                                    mySQLConsumer: String,  
                                    mySQLUser: String,  
                                    mySQLPwd: String) extends StreamingListener {  
  
  // =====  
  // Variables  
  // =====  
  
  @transient lazy val influx = InfluxDBWriter.create(influxHost)  
  @transient lazy val mysql = MySQLConnection(host = mySQLHost, table = mySQLTable, username = mySQLUser, password = mySQLPwd, database = mySQLDB)  
  @transient lazy val mySQLConnectionPool = ConnectionPool(mysql.toString).getConnection  
  
  // =====  
  // Spark Listener Override Methods (for Batches)  
  // =====  
  
  /**  
   * This method executes when a Spark Streaming batch completes.  
   *  
   * @param batchCompleted Class having information on the completed batch  
   */  
  override def onBatchCompleted(batchCompleted: StreamingListenerBatchCompleted): Unit = {  
  
    // write performance metrics to influx  
    writeBatchSchedulingStatsToInflux(batchCompleted)  
  
    // write offsets (state) to mysql  
    writeBatchOffsetsAndCounts(batchCompleted)  
  }  
}
```

# Streaming Listeners: `onBatchCompleted()`

*What data is available to us from Spark?*

- **Batch Info**
  - numRecords (total)
  - Processing times
  - delays
- **Stream Info**
  - Topic object
    - Offsets
    - numRecords (topic level)

```
batch.batchInfo.|
  SpotXSparkStreamingListener.super (tv.spotx.scala.monitoring.listeners)      SpotXSparkStreamingListener
    v batchTime
      v outputOperationInfos
      v processingEndTime
      v processingStartTime
      v streamIdToInputInfo
      v submissionTime
      m numRecords
      m processingDelay
      m schedulingDelay
      m totalDelay
    Time
      Map[Int, OutputOperationInfo]
      Option[Long]
      Option[Long]
      Map[Int, StreamInputInfo]
      Long
      Long
      Option[Long]
      Option[Long]
      Option[Long]
```

The code block shows a snippet of Scala code for a `batch.batchInfo` object. It includes a call to `super` and a list of fields. The `batchTime` field is highlighted with a blue selection bar. A tooltip or dropdown menu lists the following fields and their types:

- `outputOperationInfos`: `Map[Int, OutputOperationInfo]`
- `processingEndTime`: `Option[Long]`
- `processingStartTime`: `Option[Long]`
- `streamIdToInputInfo`: `Map[Int, StreamInputInfo]`
- `submissionTime`: `Long`
- `numRecords`: `Long`
- `processingDelay`: `Option[Long]`
- `schedulingDelay`: `Option[Long]`
- `totalDelay`: `Option[Long]`

# Streaming Listeners: *onBatchCompleted()*

We want to write performance stats to Influx using our custom method:  
***writeBatchSchedulingStatsToInflux()***

We want to write our Kafka Offsets to MySQL using our custom method:  
***writeBatchOffsetsAndCounts()***

```
/**  
 * This method executes when a Spark Streaming batch completes.  
 *  
 * @param batchCompleted Class having information on the completed batch  
 */  
override def onBatchCompleted(batchCompleted: StreamingListenerBatchCompleted): Unit = {  
  
    // write performance metrics to influx  
    writeBatchSchedulingStatsToInflux(batchCompleted)  
  
    // write offsets (state) to mysql  
    writeBatchOffsetsAndCounts(batchCompleted)  
}
```

# *writeBatchSchedulingStatsToInflux()*

This custom method consumes the Batch Info stored in the object...

- Batch duration
- Batch delay
- total records
  - across all streams/topics

... and writes them to Influx!

```
/**  
 * Pulls, parses, and logs the key performance metrics of the Streaming app and logs them to Influx.  
 * Processing Time: How many seconds needed to complete this batch (i.e. duration).  
 * Scheduling Delay: How many seconds the start time of this batch was delayed.  
 * Num Records: The total number of input records from a live stream consumed this batch.  
 *  
 * @param batch Class having information on the completed batch  
 */  
def writeBatchSchedulingStatsToInflux(batch: StreamingListenerBatchCompleted): Unit = {  
  
    // Store the processing time for this batch in seconds  
    val processingTime = if (batch.batchInfo.processingDelay.isDefined) {  
        batch.batchInfo.processingDelay.get / 1000  
    }  
    else {  
        0  
    }  
  
    // Store the scheduling delay for this batch in seconds  
    val schedulingDelay = if (batch.batchInfo.schedulingDelay.isDefined && batch.batchInfo.schedulingDelay.get > 0) {  
        batch.batchInfo.schedulingDelay.get / 1000  
    }  
    else {  
        0  
    }  
  
    // Store the total record count for this batch  
    val numRecords = batch.batchInfo.numRecords  
  
    // Log all three (3) metrics to Influx  
    influx.write(influxDB, influxMeasurement, Seq(), Seq(("processingTime", processingTime)))  
    influx.write(influxDB, influxMeasurement, Seq(), Seq(("schedulingDelay", schedulingDelay)))  
    influx.write(influxDB, influxMeasurement, Seq(), Seq(("numRecords", numRecords)))  
}
```

# *writeBatchSchedulingStatsToInflux()*



*Grafana dashboard reading from an Influx time-series database.*

# Streaming Listeners: `onBatchCompleted()`

We want to write performance stats to Influx using our custom method:

**`writeBatchSchedulingStatsToInflux()`**

We want to write our Kafka Offsets to MySQL using our custom method:  
**`writeBatchOffsetsAndCounts()`**

```
/**  
 * This method executes when a Spark Streaming batch completes.  
 *  
 * @param batchCompleted Class having information on the completed batch  
 */  
override def onBatchCompleted(batchCompleted: StreamingListenerBatchCompleted): Unit = {  
  
    // write performance metrics to influx  
    writeBatchSchedulingStatsToInflux(batchCompleted)  
  
    // write offsets (state) to mysql  
    writeBatchOffsetsAndCounts(batchCompleted)  
}
```

# *writeBatchOffsetsAndCounts()*

This custom method consumes the Stream Info stored in the object and has two steps:

- Write offsets to MySQL
- Write # records to Influx
  - per stream/topic

```
/*
 * A combination method that will handle both influx writes and MySQL offsets.
 * This is effectively a convenience method of writeBatchOffsetsToMySQL + writeBatchCountsToInflux.
 *
 * @param batch Class having information on the completed batch
 */
def writeBatchOffsetsAndCounts(batch: StreamingListenerBatchCompleted): Unit = {

    // for each stream topic consumed this batch...
    batch.batchInfo.streamIdToInputInfo.foreach(topic => {

        // write offsets for this topic to mysql
        writeTopicOffsetsToMySQL(topic)

        // write record count for this topic this batch
        writeTopicCountToInflux(topic)
    })
}
```

# *writeTopicOffsetsToMySQL()*

This custom method takes a Topic object...

... and finds the last offset processed for each kafka partition.

It then will update each partition in a database with the latest row/offset processed.

```
/***
 * Takes a topic object and writes the max offset for each partition it contains this batch to MySQL.
 *
 * @param topic A topic object within a Batch's StreamIdToInputInfo
 */
def writeTopicOffsetsToMySQL(topic: Tuple2[Int, StreamInputInfo]): Unit = {

    // map offset info to OffsetRange objects
    val partitionOffsets = topic._2.metadata("offsets").asInstanceOf[List[OffsetRange]]

    // for every partition's range of offsets
    partitionOffsets.map(offsetRange => {

        // write the new starting offset for each partition in the topic to the state db
        var maxOffset = offsetRange.untilOffset - 1

        // create a now() timestamp
        val now = new DateTime().toString("YYYY-MM-dd HH:mm:ss")

        // form the sql
        val sql =
            s"""INSERT INTO $mySQLDB.$mySQLTable (consumer, topic, partition_id, offset, offset_ts, batch_size)
               VALUES
               ('$mySQLConsumer', "${offsetRange.topic}", ${offsetRange.partition}, '$maxOffset', '$now', ${offsetRange.count})
               ON DUPLICATE KEY UPDATE offset_ts = VALUES(offset_ts), offset = VALUES(offset),
               batch_size = VALUES(batch_size)
               """
        .....

        // execute the sql to offload offsets to the table
        val st = mySQLConnectionPool.createStatement
        st.execute(sql)
        st.close()
    })
}
```

# *writeTopicOffsetsToMySQL()*

```
mysql> select * from offsets where app='myApp1' limit 10;
```

app	topic	partition_id	offset	offset_ts	batch_size	modified_date
myApp1	kafka.some_topic1	0	877744060	2019-03-27 23:30:55	836	2019-03-27 23:30:55
myApp1	kafka.some_topic1	1	877194072	2019-03-27 23:30:55	836	2019-03-27 23:30:55
myApp1	kafka.some_topic1	2	877804618	2019-03-27 23:30:55	810	2019-03-27 23:30:55
myApp1	kafka.some_topic1	3	877733267	2019-03-27 23:30:55	785	2019-03-27 23:30:55
myApp1	kafka.some_topic1	4	878811924	2019-03-27 23:30:55	859	2019-03-27 23:30:55
myApp1	kafka.some_topic1	5	877818167	2019-03-27 23:30:55	813	2019-03-27 23:30:55

*MySQL table maintaining latest offsets for an app.  
Data is at an app / topic / partition level.*

# Reading Offsets from MySQL

Your offsets are now stored in a DB after each batch completes.

Whenever your app restarts, it reads those offsets from the DB...

And starts processing where it last left off!

```
def subscribeToTopicWithState(  
    streamingContext: StreamingContext,  
    topic: String,  
    kafkaParams: Map[String, Object],  
    dbConnection: MySQLConnection,  
    stateConsumer: String,  
    resetOffsets: Boolean = false  
): InputDStream[ConsumerRecord[String, Option[GenericRecord]]] = {  
  
    // Store list of topics in an array  
    val topics = Array(topic)  
  
    // Either Get Offsets from a State Database or Reset to Latest  
    val offsets = getOffsetsFromState(dbConnection, stateConsumer, topic, resetOffsets)  
  
    // Subscribe to the topic based on the offsets provided  
    val strategy = getStrategy(offsets, kafkaParams, topics)  
  
    // Create a direct stream and pull data  
    val subscribed = KafkaUtils.createDirectStream[String, Option[GenericRecord]](  
        streamingContext,  
        LocationStrategies.PreferConsistent,  
        strategy  
    )  
  
    // Filter data coming out of Kafka to remove bad records  
    subscribed  
}
```

# Getting Offsets from the Database

```
def getOffsetsFromState(dbConnection: MySQLConnection,
                      consumer: String,
                      topic: String,
                      resetOffsets: Boolean
                     ): Map[TopicPartition, Long] = {

    // Establish a MySQL Connection given a URL
    val stateDbConn = ConnectionPool(dbConnection.toString).getConnection

    // If You want to reset offsets, just return an empty map
    val offsets =
        if (resetOffsets) {
            new HashMap[TopicPartition, Long]()
        }

    // If you want to use your existing offsets, get them from MySQL
    else {
        val stmt = stateDbConn.createStatement
        val query = "SELECT * FROM " + dbConnection.database + "." + dbConnection.table +
                    " WHERE consumer = '" + consumer + "' AND topic = '" + topic + "'"
        logger.info(query)
        val result = stmt.executeQuery(query)
        new RsIterator(result).map(
            x =>
                new TopicPartition(
                    x.getString("topic"),
                    x.getInt("partition_id")
                ) ->
                    (x.getLong("offset") + 1L) // Add one because we record last known offset
            ).toMap
    }

    offsets
}
```

# Example: Reading Offsets from MySQL

```
def subscribeToTopicWithState(
    streamingContext: StreamingContext,
    topic: String,
    kafkaParams: Map[String, Object],
    dbConnection: MySQLConnection,
    stateConsumer: String,
    resetOffsets: Boolean = false
): InputDStream[ConsumerRecord[String, Option[GenericRecord]]] = {

    // Store list of topics in an array
    val topics = Array(topic)

    // Either Get Offsets from a State Database or Reset to Latest
    val offsets = getOffsetsFromState(dbConnection, stateConsumer, topic, resetOffsets)

    // Subscribe to the topic based on the offsets provided
    val strategy = getStrategy(offsets, kafkaParams, topics)

    // Create a direct stream and pull data
    val subscribed = KafkaUtils.createDirectStream[String, Option[GenericRecord]](
        streamingContext,
        LocationStrategies.PreferConsistent,
        strategy
    )

    // Filter data coming out of Kafka to remove bad records
    subscribed
}
```

# Example: Reading Offsets from MySQL

```
def getStrategy(
    offsets: Map[TopicPartition, Long],
    kafkaParams: Map[String, Object],
    topics: Array[String]
): ConsumerStrategy[String, Option[GenericRecord]] ={

    // if offsets were pulled from MySQL, assign the stream to start there
    val strategy = if (offsets.size > 0) {
        ConsumerStrategies.Assign[String, Option[GenericRecord]](offsets.keys.toList, kafkaParams, offsets)
    }

    // otherwise, start at latest data
    else {
        ConsumerStrategies.Subscribe[String, Option[GenericRecord]](topics, kafkaParams)
    }

    strategy
}
```

# Reading Offsets from MySQL

```
def subscribeToTopicWithState(
    streamingContext: StreamingContext,
    topic: String,
    kafkaParams: Map[String, Object],
    dbConnection: MySQLConnection,
    stateConsumer: String,
    resetOffsets: Boolean = false
): InputDStream[ConsumerRecord[String, Option[GenericRecord]]] = {

    // Store list of topics in an array
    val topics = Array(topic)

    // Either Get Offsets from a State Database or Reset to Latest
    val offsets = getOffsetsFromState(dbConnection, stateConsumer, topic, resetOffsets)

    // Subscribe to the topic based on the offsets provided
    val strategy = getStrategy(offsets, kafkaParams, topics)

    // Create a direct stream and pull data
    val subscribed = KafkaUtils.createDirectStream[String, Option[GenericRecord]](
        streamingContext,
        LocationStrategies.PreferConsistent,
        strategy
    )

    // Filter data coming out of Kafka to remove bad records
    subscribed
}
```

# *writeBatchOffsetsAndCounts()*

This custom method has two steps:

- Write offsets to MySQL
- Write # records to Influx
  - per stream/topic

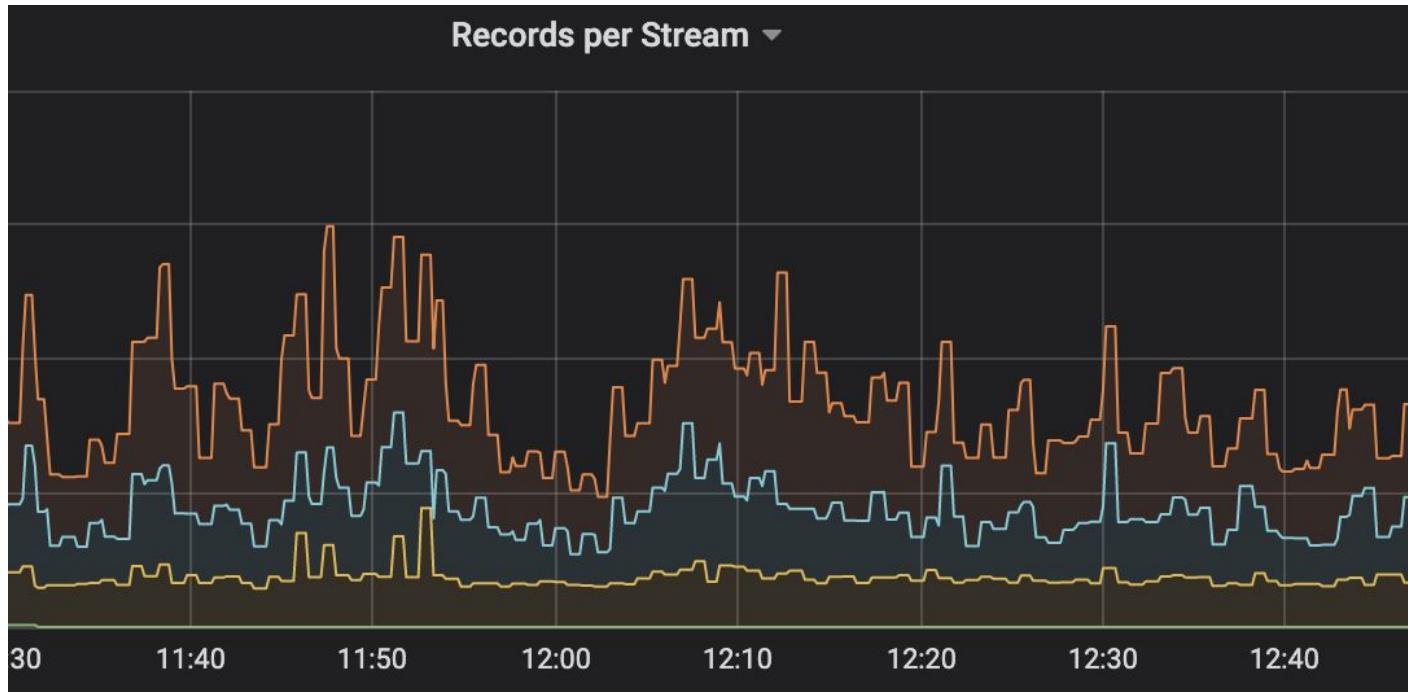
```
/**  
 * A combination method that will handle both influx writes and MySQL offsets.  
 * This is effectively a convenience method of writeBatchOffsetsToMySQL + writeBatchCountsToInflux.  
 *  
 * @param batch Class having information on the completed batch  
 */  
def writeBatchOffsetsAndCounts(batch: StreamingListenerBatchCompleted): Unit = {  
  
    // for each stream topic consumed this batch...  
    batch.batchInfo.streamIdToInputInfo.foreach(topic => {  
  
        // write offsets for this topic to mysql  
        writeTopicOffsetsToMySQL(topic)  
  
        // write record count for this topic this batch  
        writeTopicCountToInflux(topic)  
    }  
}
```

# *writeTopicCountToInflux()*

This method takes a *Topic* object and writes the # of records processed for a single topic (more granular than batch-level record total).

```
/**  
 * Takes a topic object and writes the number of records for said topic this batch to Influx.  
 *  
 * @param topic A topic object within a Batch's StreamIdToInputInfo  
 */  
def writeTopicCountToInflux(topic: Tuple2[Int, StreamInputInfo]): Unit = {  
  
    // store the individual record count for this topic  
    val numRecords = topic._2.numRecords  
  
    // store topicName  
    val topicName = topic._2.metadata("offsets").asInstanceOf[List[OffsetRange]].head.topic  
  
    // write record count for this topic this batch  
    influx.write(influxDB, influxMeasurement, Seq(), Seq(("numRecords_" + topicName, numRecords)))  
}
```

# *writeTopicCountToInflux()*



*Grafana dashboard reading from an Influx time-series database.*

# Streaming Listeners: Add to Context

Finally, instantiate your listener within an application.

You can customize with any args you need. By default it needs none.

Use the *addStreamingListener()* method of your Spark Streaming Context! Done!

```
val sparkConf = new SparkConf()
    .setAppName("My Spark App")
    .setMaster("local[*]")

val ssc = new StreamingContext(sparkConf, Seconds(30))

val listener = new SpotXSparkStreamingListener(
    "influx.yourcompany.com",
    "my_influx_database",
    "my_measurement",
    "myrdbms.yourcompany.com",
    "my_rdbms_database",
    "my_rdbms_table",
    "my_kafka_group_id",
    "my_rdbms_user",
    "my_rdbms_pwd")

ssc.addStreamingListener(listener)
```

# The Possibilities are numerous!

The *Developer API* has many other data points you can build great monitoring from, such as:

- StreamingJobProgressListener
  - waitingBatches
  - runningBatches
  - numReceivers
  - lastCompletedBatch
  - Just to name a few!

Method and Description
batchDuration()
lastCompletedBatch()
lastReceivedBatch()
lastReceivedBatchRecords()
numReceivers()
numTotalCompletedBatches()
numTotalProcessedRecords()
numTotalReceivedRecords()
numUnprocessedBatches()
onBatchCompleted(StreamingListenerBatchCompleted batchCompleted)
Called when processing of a batch of jobs has completed.
onBatchStarted(StreamingListenerBatchStarted batchStarted)
Called when processing of a batch of jobs has started.
onBatchSubmitted(StreamingListenerBatchSubmitted batchSubmitted)
Called when a batch of jobs has been submitted for processing.
onReceiverError(StreamingListenerReceiverError receiverError)
Called when a receiver has reported an error
onReceiverStarted(StreamingListenerReceiverStarted receiverStarted)
Called when a receiver has been started
onReceiverStopped(StreamingListenerReceiverStopped receiverStopped)
Called when a receiver has been stopped
processingDelayDistribution()
receivedRecordsDistributions()
receiverInfo(int receiverId)
retainedCompletedBatches()
runningBatches()
schedulingDelayDistribution()
totalDelayDistribution()
waitingBatches()

# A Word About Batch Listeners

Allow the extension of non-streaming events and triggers for each application, job, stage, and executor.

Valuable usage:

- Write statistics for runtimes of applications (given appId or appName, both available)
- Write info about executor removal
  - A reason is provided
- Write each Job duration, identifying long running jobs
  - Find potential skew
  - Find potential problem nodes

Executors					
	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks
Active(124)	248	4.4 MB / 4.2 GB	0.0 B	492	209
Dead(0)	0	0.0 B / 0.0 B	0.0 B	0	0
Total(124)	248	4.4 MB / 4.2 GB	0.0 B	492	209

Active Stages (1)						
Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	
0	sql at HiveOrcConverter.scala:41 +details (kill)	2019/03/28 20:30:24	8.2 min	3326/6577	65.6 GB	

# Batch Listeners: *SparkListener* Abstract Class

Similar to streaming, everything you'll need is in the *SparkListener*, but this time in an *Abstract Class* within the Developer API.

However, it has 18 convenience methods for overriding.

We find  
*onApplicationStart()*,  
*onApplicationEnd()*, and  
*onExecutorRemoved()*  
to be the most useful.

```
/**  
 * :: DeveloperApi ::  
 * A default implementation for [[SparkListenerInterface]] that has no-op implementations for  
 * all callbacks.  
 *  
 * Note that this is an internal interface which might change in different Spark releases.  
 */  
@DeveloperApi  
abstract class SparkListener extends SparkListenerInterface {  
    override def onStageCompleted(stageCompleted: SparkListenerStageCompleted): Unit = {}  
    override def onStageSubmitted(stageSubmitted: SparkListenerStageSubmitted): Unit = {}  
    override def onTaskStart(taskStart: SparkListenerTaskStart): Unit = {}  
    override def onTaskGettingResult(taskGettingResult: SparkListenerTaskGettingResult): Unit = {}  
    override def onTaskEnd(taskEnd: SparkListenerTaskEnd): Unit = {}  
    override def onJobStart(jobStart: SparkListenerJobStart): Unit = {}  
    override def onJobEnd(jobEnd: SparkListenerJobEnd): Unit = {}  
    override def onEnvironmentUpdate(environmentUpdate: SparkListenerEnvironmentUpdate): Unit = {}  
    override def onBlockManagerAdded(blockManagerAdded: SparkListenerBlockManagerAdded): Unit = {}  
    override def onBlockManagerRemoved(  
        blockManagerRemoved: SparkListenerBlockManagerRemoved): Unit = {}  
    override def onUnpersistRDD(unpersistRDD: SparkListenerUnpersistRDD): Unit = {}  
    override def onApplicationStart(applicationStart: SparkListenerApplicationStart): Unit = {}  
    override def onApplicationEnd(applicationEnd: SparkListenerApplicationEnd): Unit = {}  
    override def onExecutorMetricsUpdate(  
        executorMetricsUpdate: SparkListenerExecutorMetricsUpdate): Unit = {}  
    override def onExecutorAdded(executorAdded: SparkListenerExecutorAdded): Unit = {}  
    override def onExecutorRemoved(executorRemoved: SparkListenerExecutorRemoved): Unit = {}  
    override def onBlockUpdated(blockUpdated: SparkListenerBlockUpdated): Unit = {}  
    override def onOtherEvent(event: SparkListenerEvent): Unit = {}  
}
```

# Batch Examples

```
class BatchListenerUtils(  
    measurementName: String,  
    database: String) extends SparkListener {  
  
    override def onApplicationStart(applicationStart: SparkListenerApplicationStart): Unit = {  
        influxWriter.write(myTimeSeriesTS,  
                           applicationStart.appName,  
                           applicationStart.time)  
    }  
  
    override def onExecutorRemoved(executorRemoved: SparkListenerExecutorRemoved): Unit = {  
        influxWriter.write(myTimeSeriesTS,  
                           executorRemoved.reason)  
    }  
}
```

# Review

Monitoring spark apps is tough!

**But** we found a solution: **Spark Listeners!**

Streaming  
Batch

Both provide actionable data:

- Performance metrics
- Data volume





SPARK+AI  
SUMMIT 2019

# Q & A

#UnifiedAnalytics #SparkAISummit

# Code!

Blog: [hadoopsters.dev](https://hadoopsters.dev)

Github: [gist.github.com/hadoopsters](https://gist.github.com/hadoopsters)

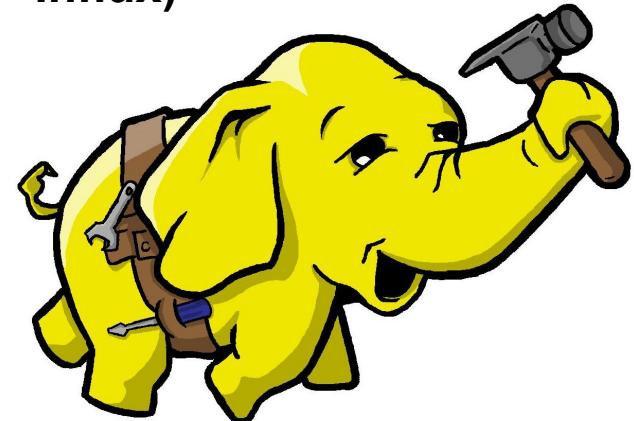
- [ExampleStreamingListener.scala](#) (Generic implementation)
- [SpotXSparkStreamingListener.scala](#) (MySQL + Influx)

Landon Robinson

- [\*\*lrobinson@spotx.tv\*\*](mailto:lrobinson@spotx.tv)

Ben Storrie

- [\*\*bstorrie@spotx.tv\*\*](mailto:bstorrie@spotx.tv)





SPARK+AI  
SUMMIT 2019

DON'T FORGET TO RATE  
AND REVIEW THE SESSIONS

SEARCH SPARK + AI SUMMIT



Download on the  
App Store



GET IT ON  
Google Play

