

**REAL-TIME BIG DATA ANALYTICS
IN THE CLOUD 101: EXPERT
ADVICE FROM THE ATTUNITY AND
AZURE DATA LAKE STORAGE
GEN2 TEAMS**



Today's Speakers



Carole Gunst
Marketing Director
Attunity



Jordan Martz,
Director Technology Solutions,
Attunity



Jeff King,
Senior Program Manager for Azure Data Lake Storage Gen2,
Microsoft

We'll answer these questions

Why is real-time data important for driving business insights?

What's a data lake and why would you use one to store your real-time data?

How can you use change data capture (CDC) technology to efficiently transfer data to the cloud?

How can you build sophisticated analytic workflows quickly?

Why is Azure Data Lake Storage Gen2 the best data lake for real-time analytics?

Migrating Data to the Microsoft Data Platform

ATTUNITY


Jordan Martz, Director of Technology Solutions, Attunity

Microsoft + Attunity Strategic Partnership



Cloud and database migrations

to Microsoft Data Platform from any enterprise platform -- on-premises and in Azure



Ingests into Microsoft's DBs, EDWs and data lakes

from a broad range of data sources: SAP, Oracle, Teradata, mainframes, and more



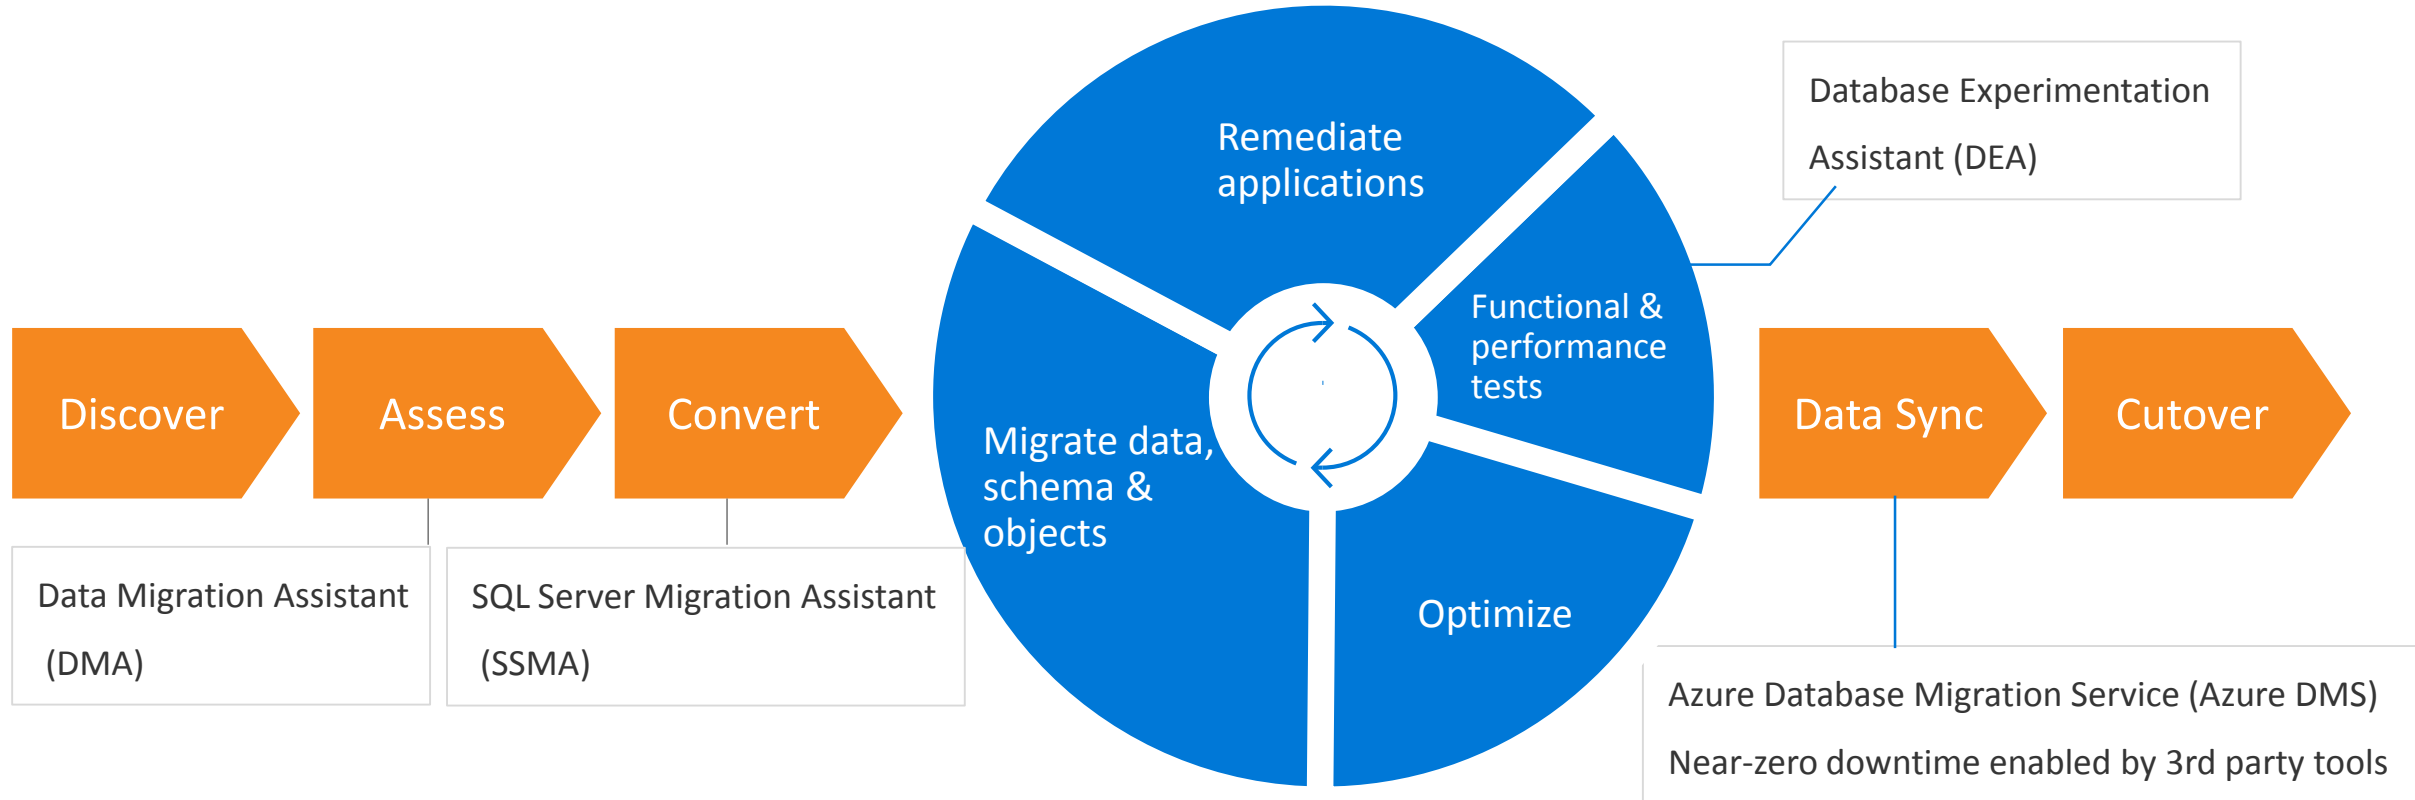
Continuous refresh of data

for zero-downtime migrations, real-time streaming for business intelligence & analytics

#1

DATA MIGRATION & CDC TECHNOLOGY

Data Migration Lifecycle









Additional Sources and Microsoft Targets Available



SOURCES

each
with CDC

 RDBMS	 DW	 HADOOP	 CLOUD	 MAINFRAME	 SAP
Oracle SQL Server DB2 iSeries DB2 z/OS DB2 LUW MySQL PostgreSQL Sybase ASE Informix	Exadata Teradata Netezza Vertica	Hortonworks Cloudera MapR	AWS RDS Salesforce	DB2 for z/OS IMS/DB VSAM SQL/MP Enscribe RMS	ECC on Oracle ECC on SQL ECC on DB2

Microsoft targets

 RDBMS	 DW	 HADOOP	 CLOUD	 STREAMING
SQL Server MySQL PostgreSQL	Microsoft PDW Azure SQL DW	Microsoft HDI	Azure SQL DW Azure SQL DB Azure Data Lake	Azure Event Hubs Kafka

Attunity Replicate – Optimized and Secure

*“We load **+10 Billion rows/hour** doing full load while replicating from Oracle Exadata with Attunity Replicate.”*

*“Attunity handles **460,000 records/sec** doing CDC from large and highly active Oracle databases into a data warehouse, with peaks of 100Gb per hour”*



PERFORMANCE-OPTIMIZED DATA TRANSFER

Extraction

optimized with
CDC and batch
bulk-loads

Transfer on-premises

optimized with
in-memory
streaming

Transfer to Azure

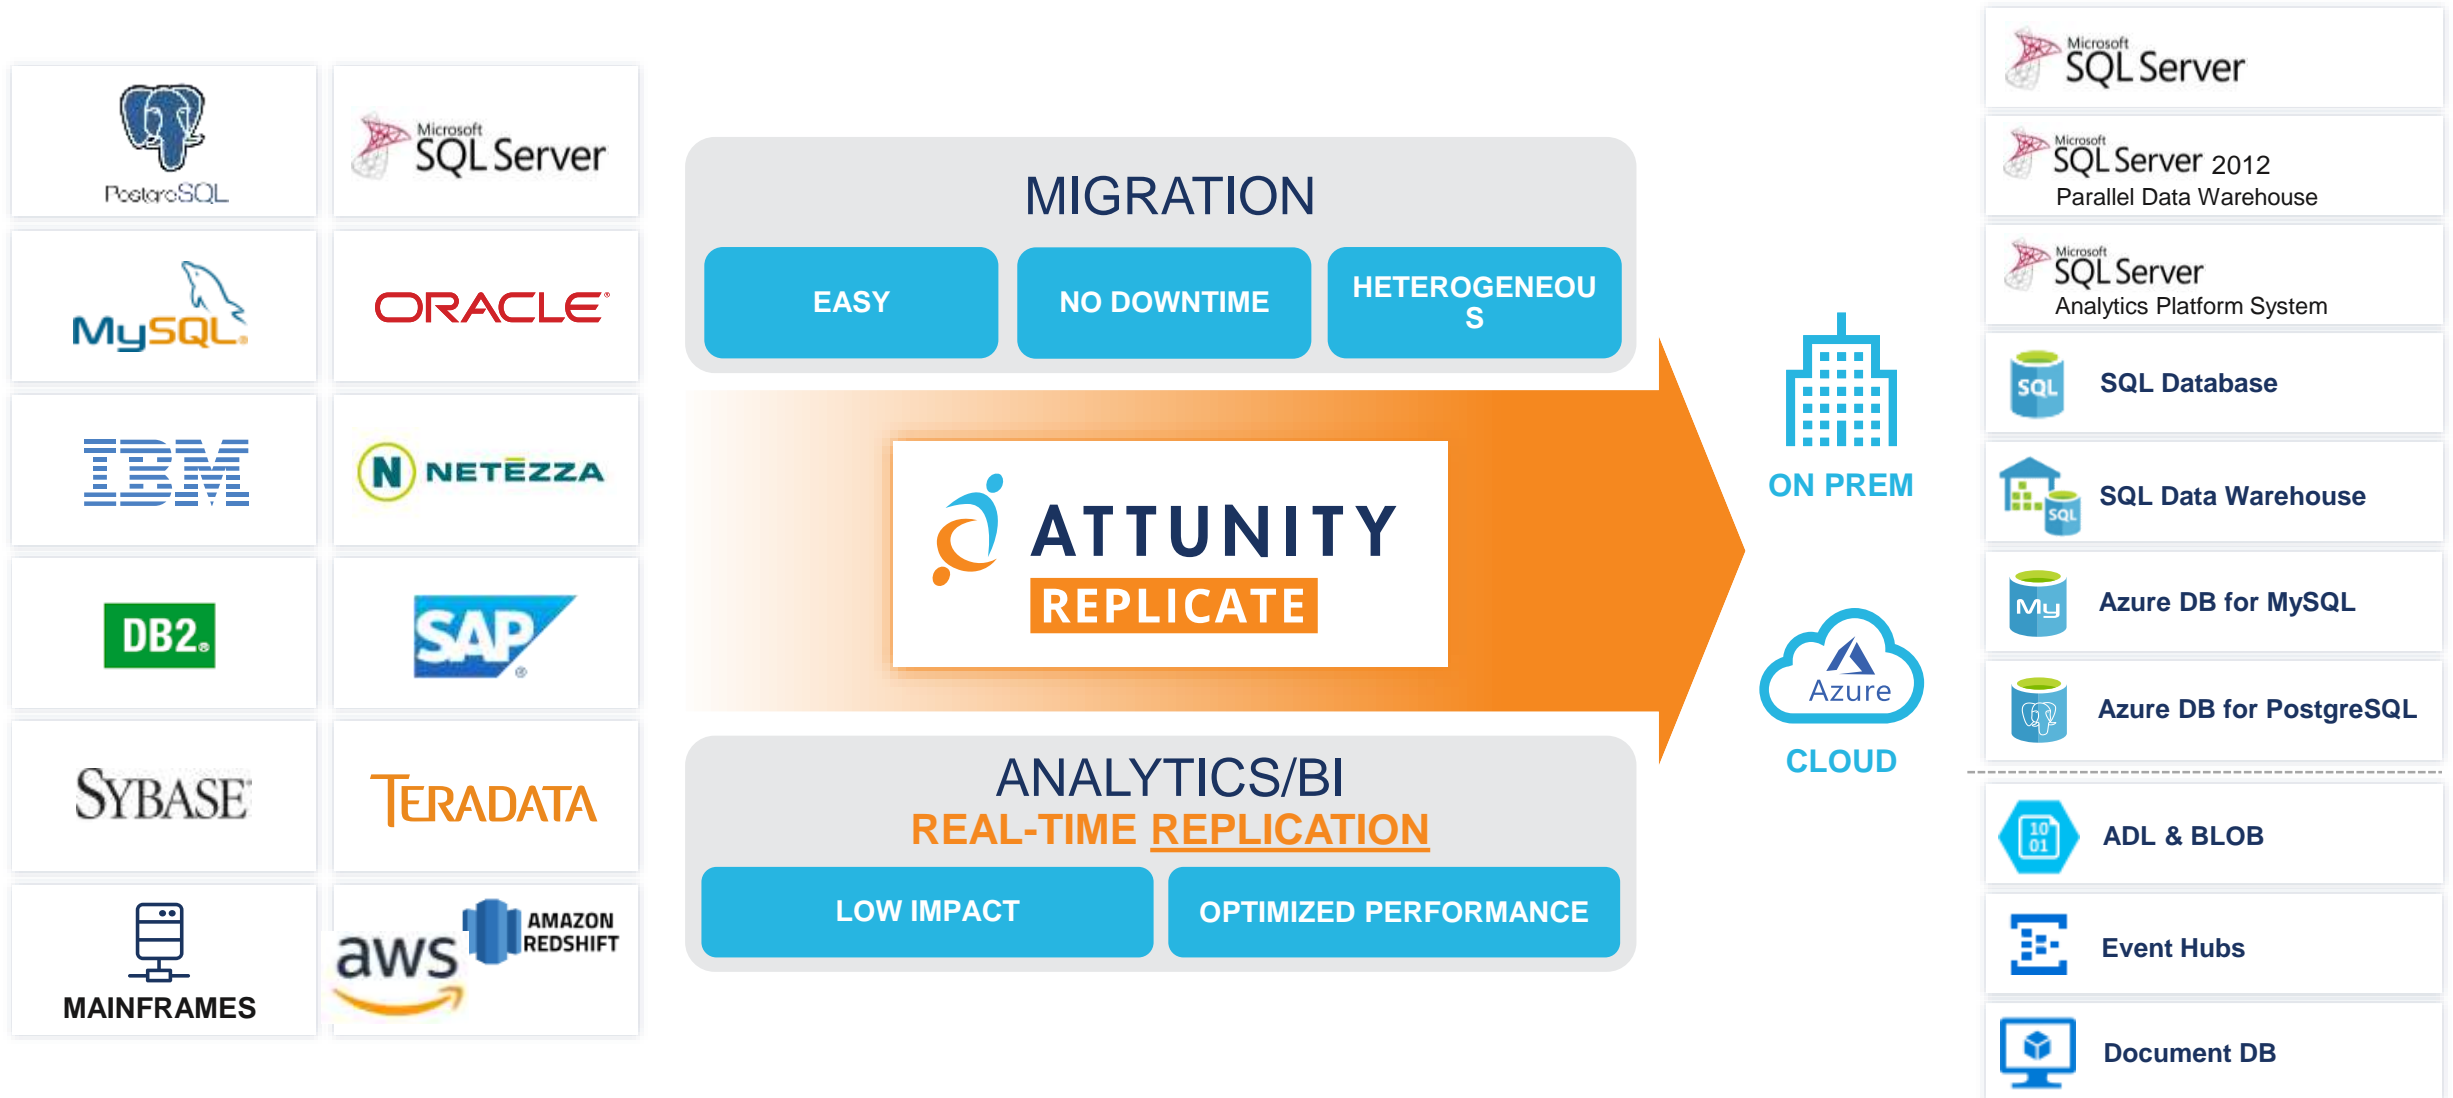
optimized file
transfer &
compression

Ingest

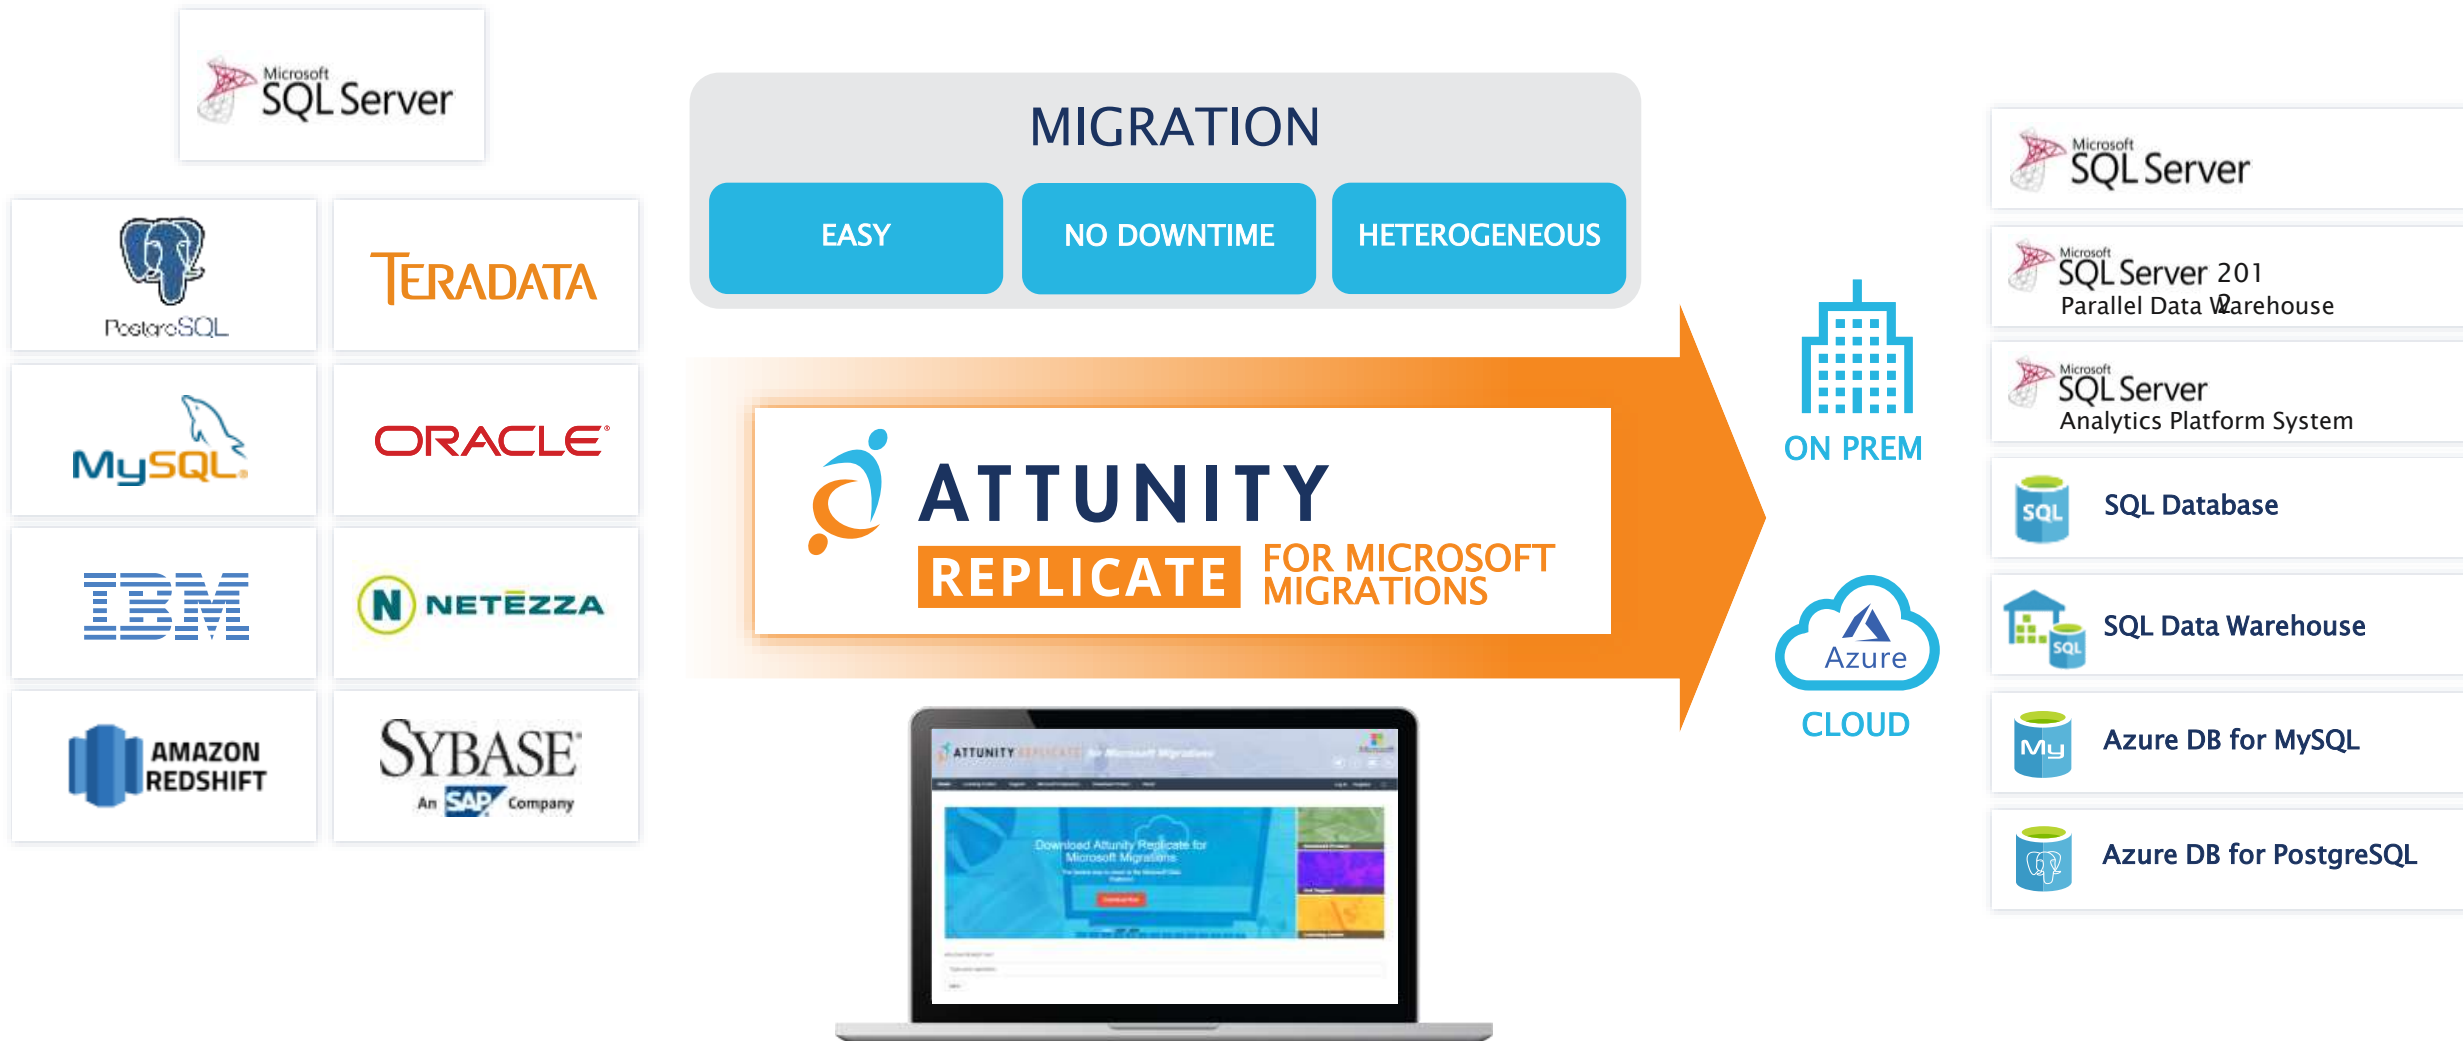
optimized for
each different
target (RDBMS,
DW, Hadoop)

SECURE

Universal Solution for the Microsoft Data Platform

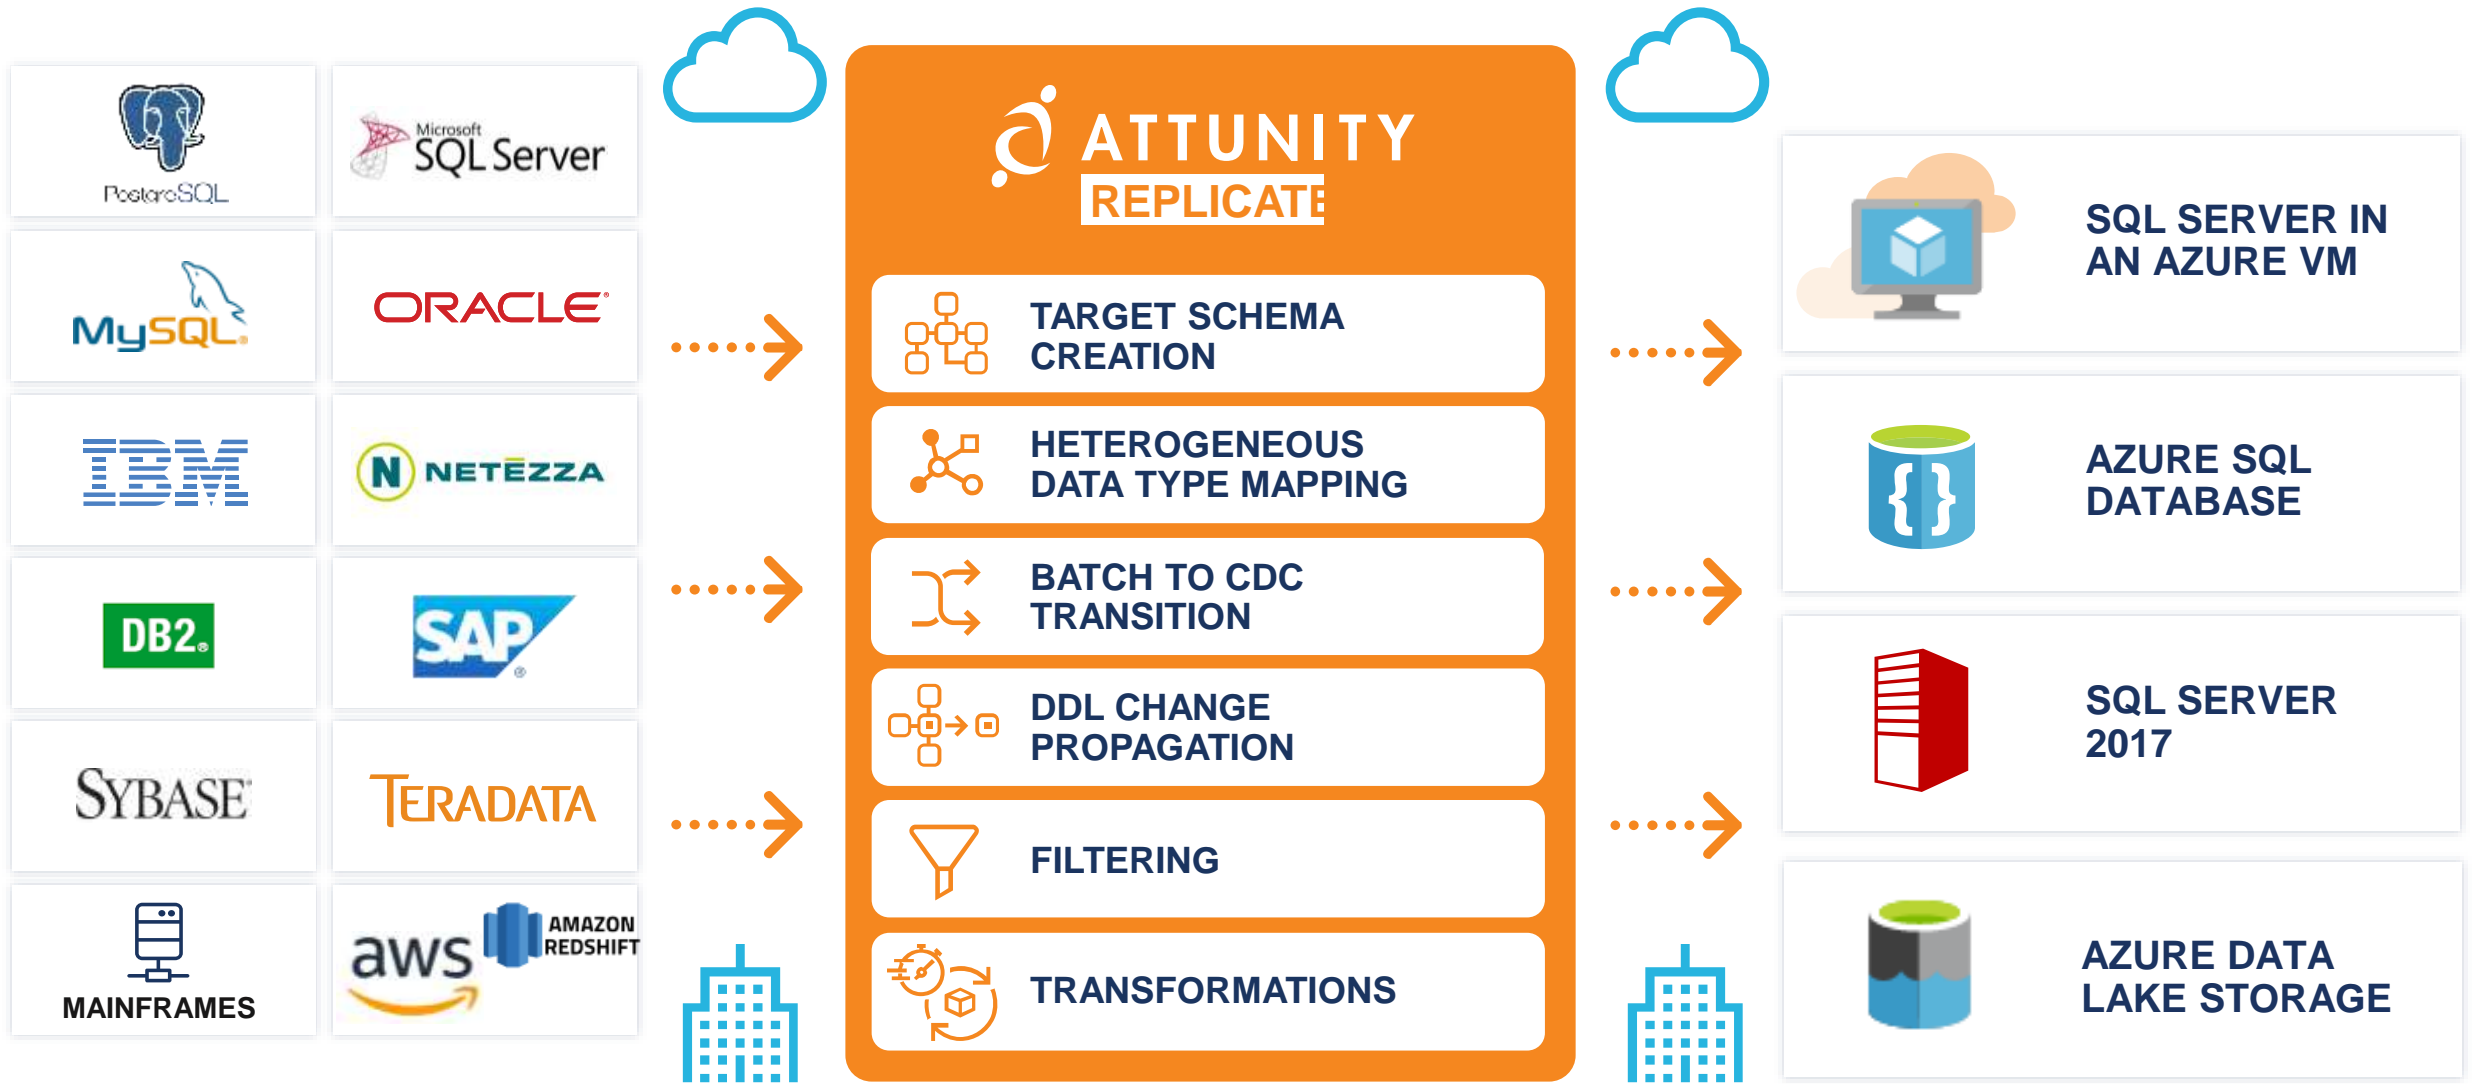


Microsoft Data Platform Migration

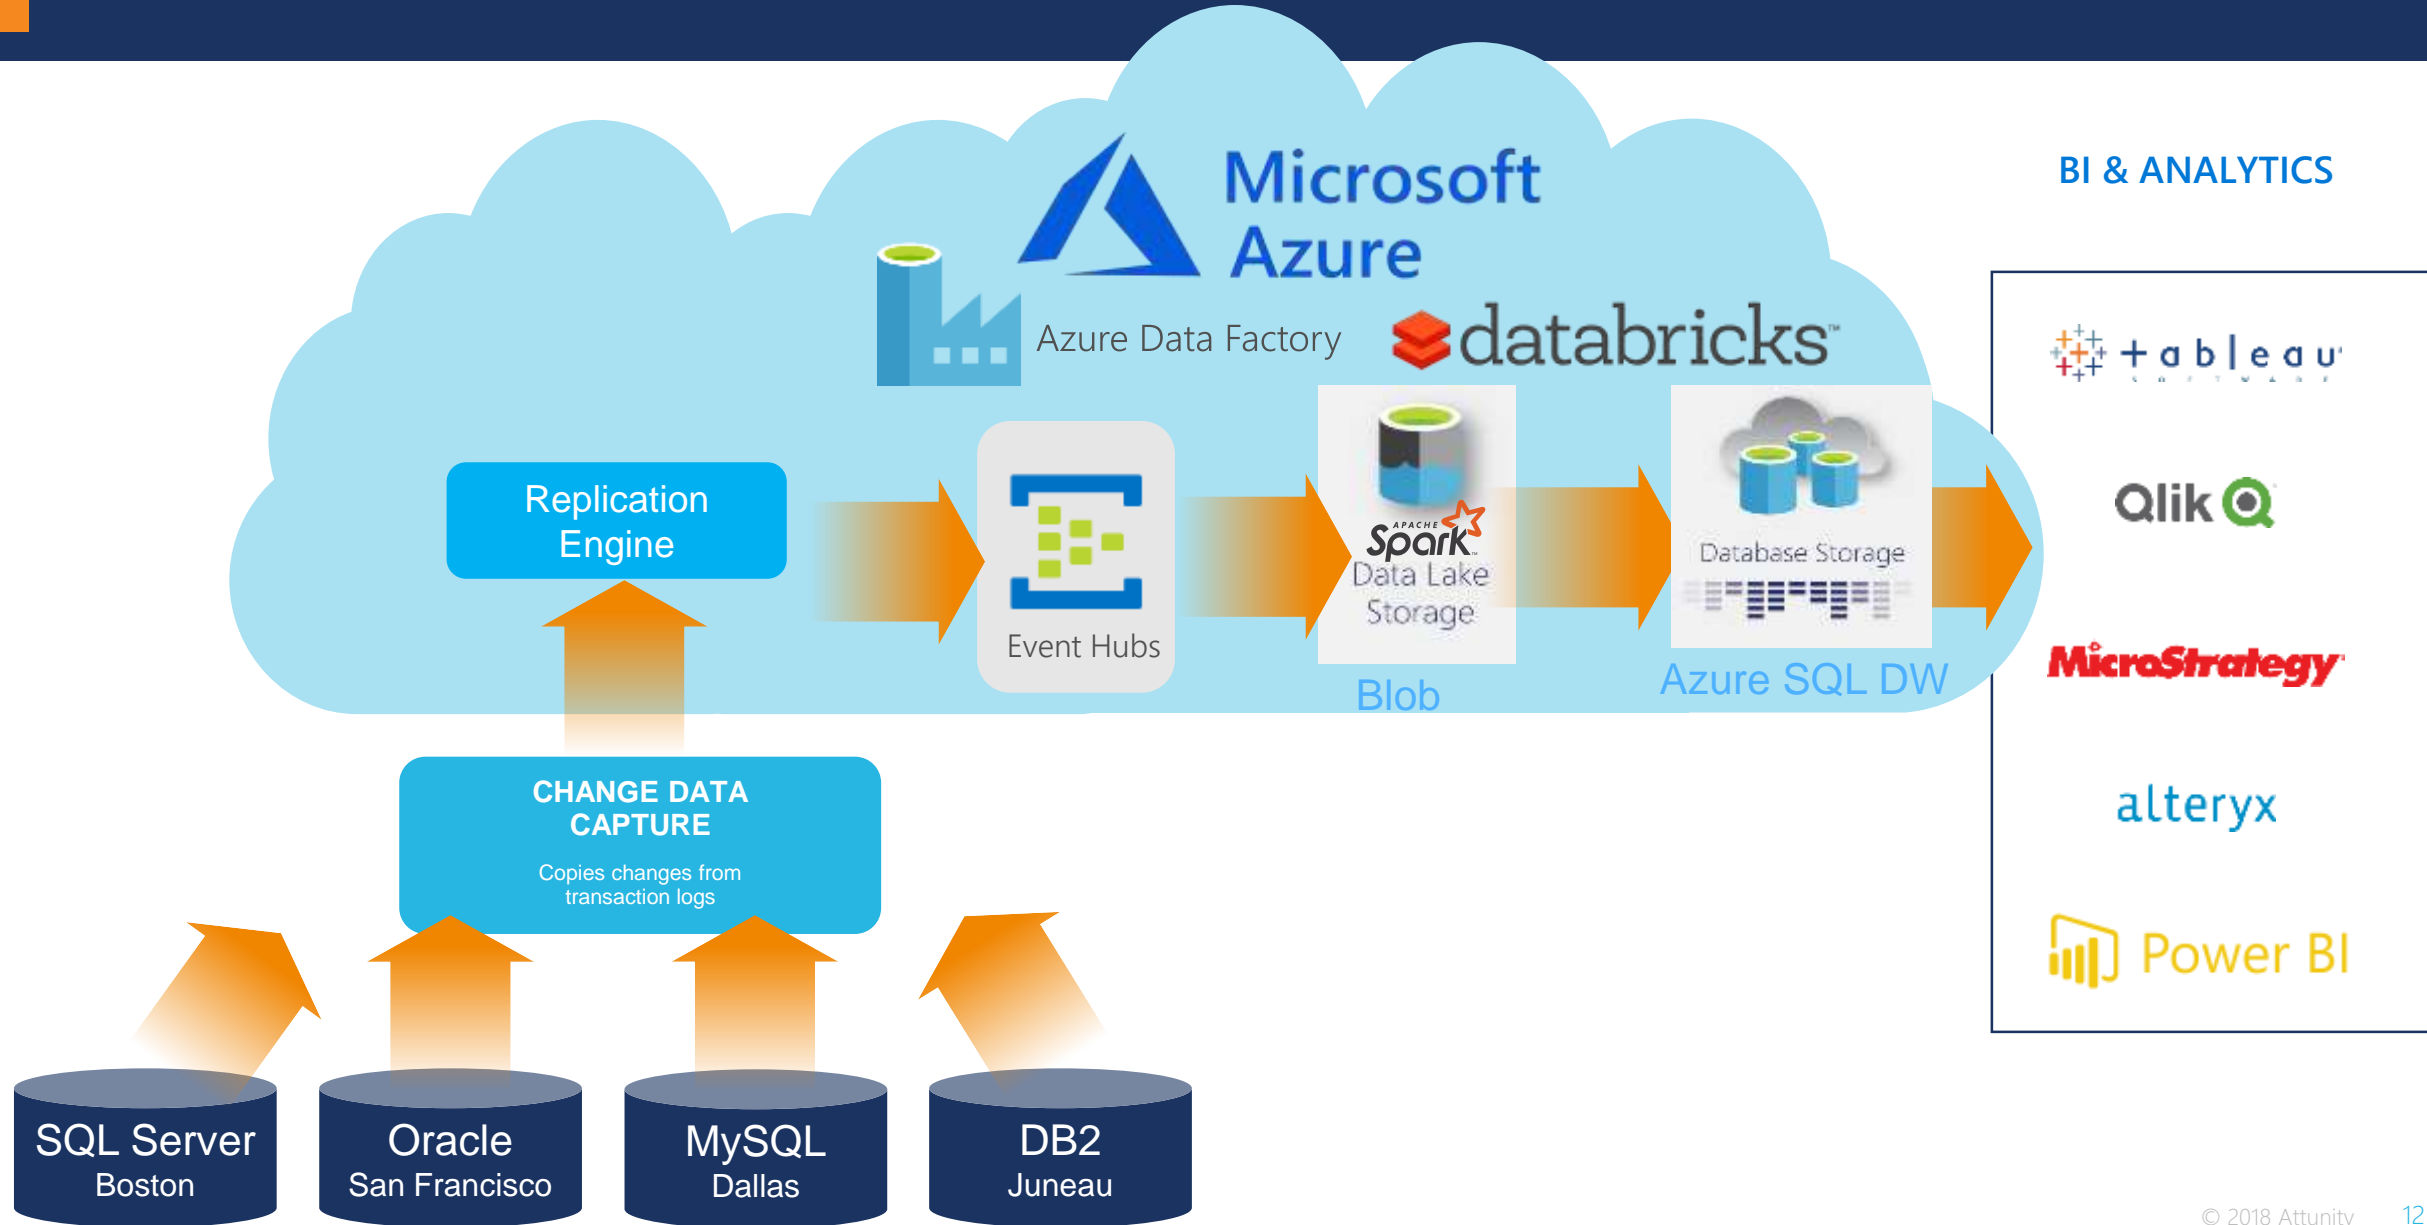


www.attunity.com/MicrosoftMigrations

Data Migration & On-Going Replication with Attunity



Architecture



Attunity Replicate Compliments Azure Data Factory



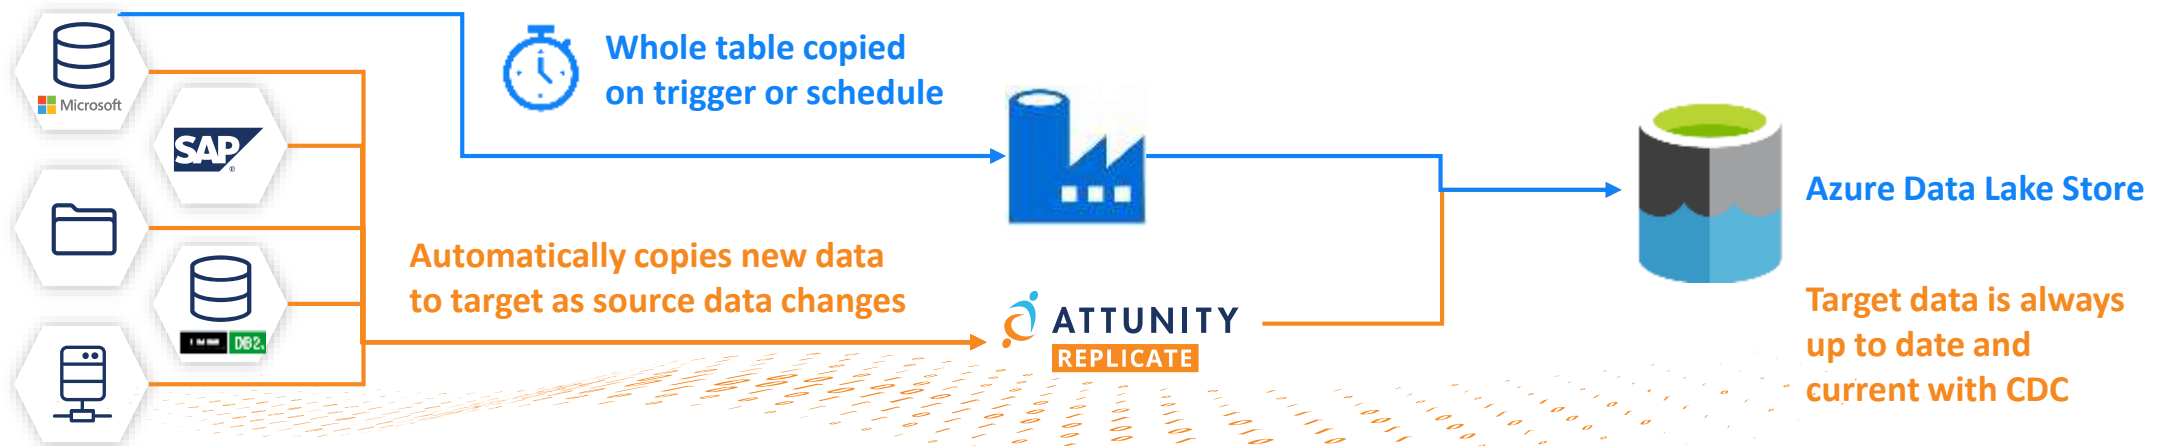
Azure Data Factory

- Data Integration as a Service via fully managed ETL workflow
- Lift-&-Shift compatibility with SSIS integration projects
- Bulk-load connectivity to many source systems
- Easy to use by developers and DBAs



- Hi-speed/Hi-volume data ingest for heterogeneous migrations & Big Data analytics
- Automated Change Data Capture (CDC), including replication of changes in metadata (DDL) and data itself
- Support for legacy and enterprise sources to Microsoft Data Services and SQL on-prem
- Easy to use by DBAs – no coding experience

Why Use Attunity Replicate If You Have Azure Data Factory?



Attunity Replicate Features	Attunity Replicate Benefits
Incremental Copy	<ul style="list-style-type: none">Copy just changed rows from any data source in a true CDC manner, not requiring delta table management
Automated Data Transfer	<ul style="list-style-type: none">No TRIGGERS or SCHEDULES – Data moves from source to target as soon as it changes on the source.
High Speed, Large Volume Data Transport	<ul style="list-style-type: none">Optimized for fast data loading into any Azure target(s)*. (*Attunity Replicate can copy source data to multiple targets)
Legacy Systems and Enterprise Applications	<ul style="list-style-type: none">Natively supports legacy, EDW and enterprise systems data sources - Mainframe, iSeries, Exadata, Teradata, Netezza, Vertica, Grenplum/Pivotal, DB2 for z/OS, IMS/DB, VSAM, SAP, SAP Hana, SQL/MP, Oracle, Microsoft SQL Server, DB2 iSeries, DB2 z/OS, DB2 LUW, MySQL, PostgreSQL, Sybase ASE, Informix, Apache Kafka, and many more.
Example Use Case (above)	<ul style="list-style-type: none">A retailer copies their ERP data from IBM DB2 to Azure Data Lake Store using Azure Data Factory.However their Point of Sales (POS) tables are 200-300GB and transfer is slow with ADF ODBC.Attunity Replicate loads the data incrementally (as it changes in real time)

Microsoft Use Cases Enabled by Attunity Solutions



SQL Server Modernization:
Upgrade from SQL 2005, 2008, etc in one step



Competitive Database Migrations:
Move from Oracle, Sybase, DB2, etc.



Competitive Data Warehouse Migrations: from Teradata, Netezza, AWS Redshift

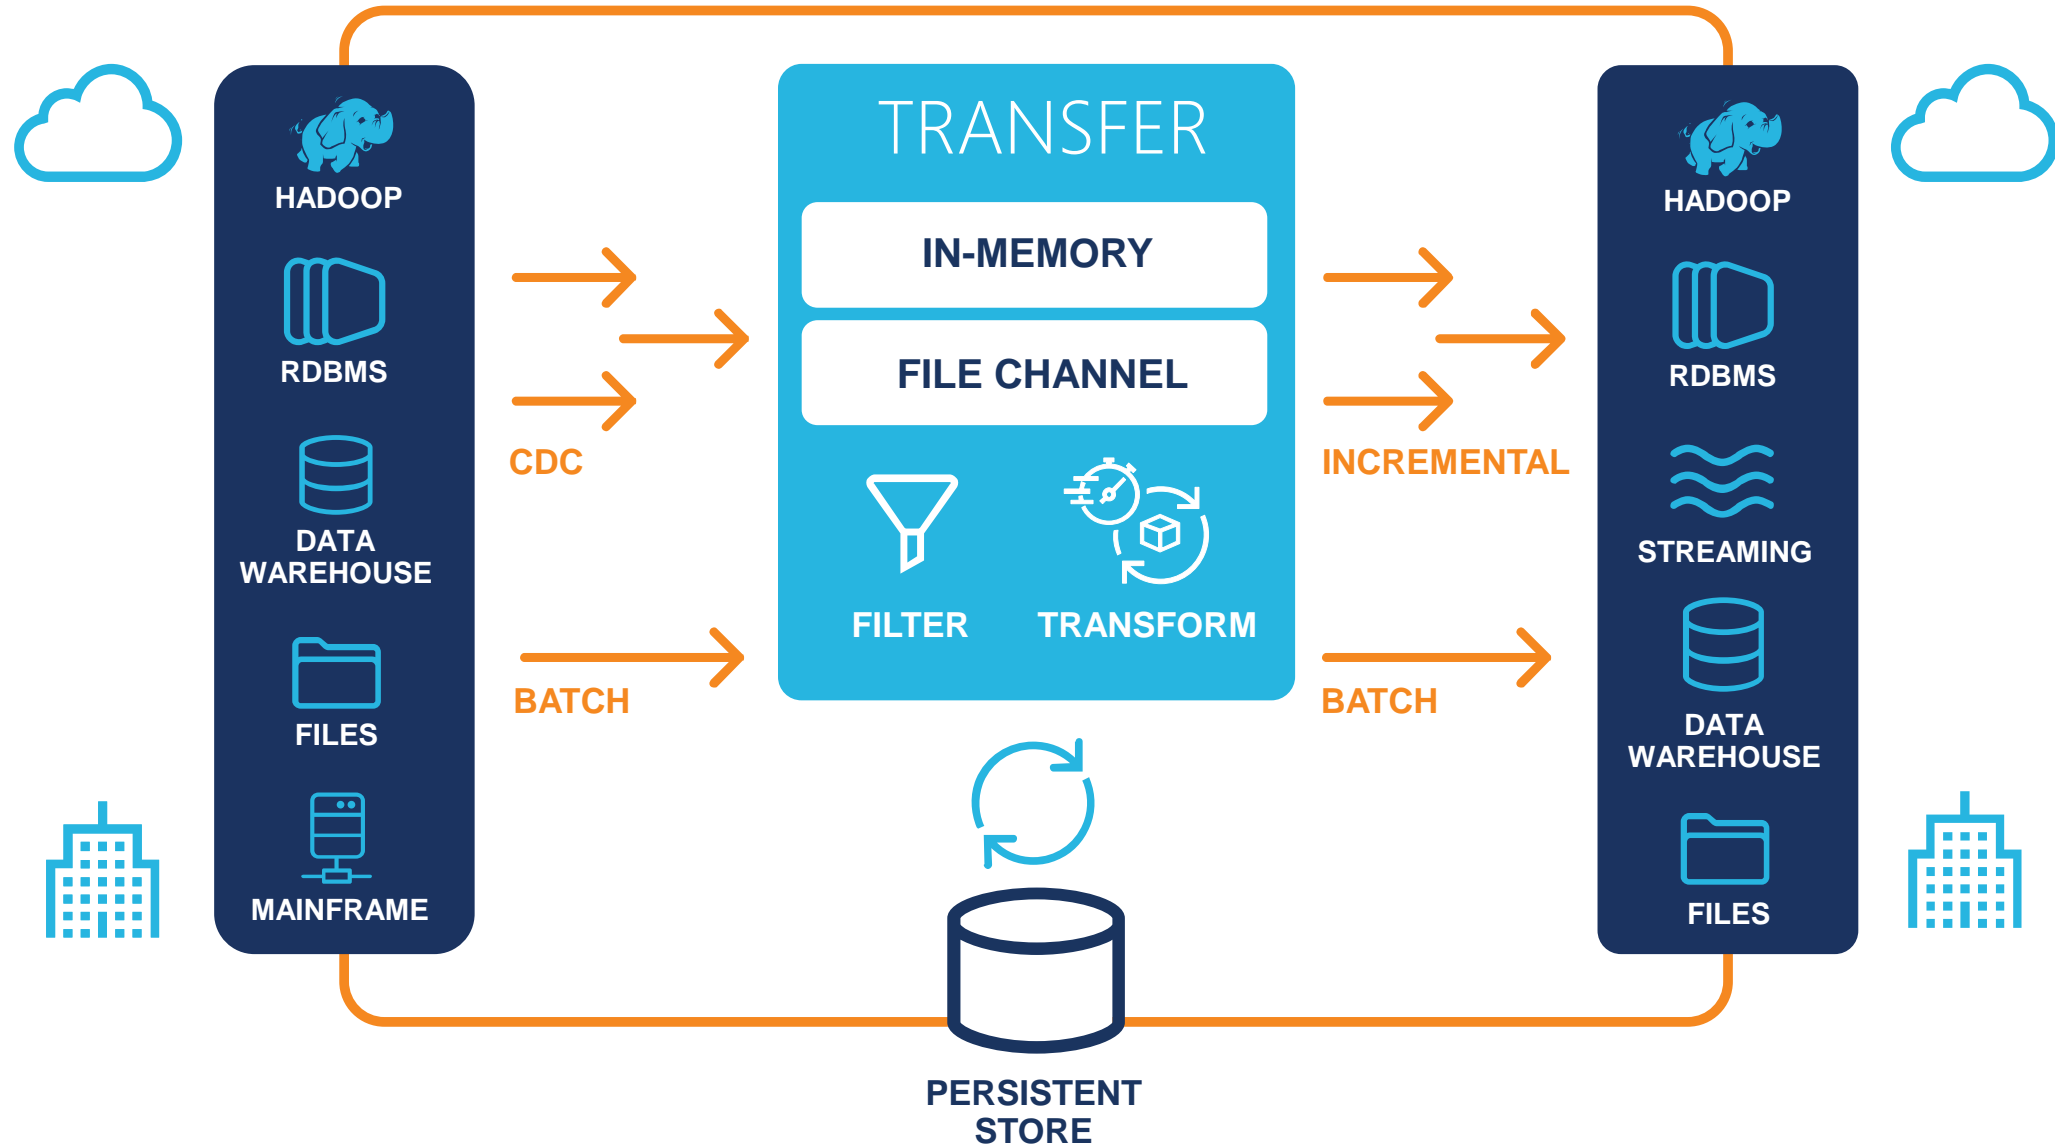


Azure Data Lake Store Gen 2 for BI / Analytics:
Ingest data from Operational Data Stores



Real-time streaming into Azure Event Hubs and HDInsight for Stream Analytics

Attunity Replicate Architecture



Azure Data Lake Storage Gen 2

MICROSOFT

Jeff King, Senior Program Manager for Azure Data Lake Storage Gen2, Microsoft

What Makes a Great Data Lake?

Massive scale

PB-scale, data
accessible
everywhere,
growth on
demand

Secure

Granular security
and protection
against accidental
data loss

Optimized for maximum performance

Lightning-
quick job
execution

Integration Friendly

Supports multiple
methods of data
ingress,
processing, egress
and visualization

Cost Effectiveness

Cloud economic
model with
ability to
intelligently
manage costs

Rich data management and governance

Azure Data Lake Storage Gen2

A **“no-compromises”** data lake: secure, performant, massively-scalable data lake storage that brings the cost and scale profile of object storage together with the performance and analytics feature set of data lake storage



SECURE

- ✓ Support for fine-grained ACLs, protecting data at the file and folder level
- ✓ Multi-layered protection via at-rest Storage Service encryption and Azure Active Directory integration



MANAGEABLE

- ✓ Automated Lifecycle Policy Management
- ✓ Object Level tiering



FAST

- ✓ Atomic file operations means jobs complete faster
- ✓ High throughput



SCALABLE

- ✓ No limits on data store size
- ✓ Global footprint (50 regions)



COST EFFECTIVE

- ✓ Object store pricing levels
- ✓ File system operations minimize transactions required for job completion



INTEGRATION READY


- ✓ Optimized for Spark and Hadoop Analytic Engines
- ✓ Tightly integrated with Azure end to end analytics solutions

A Good Data Lake Should be Multi-modal...

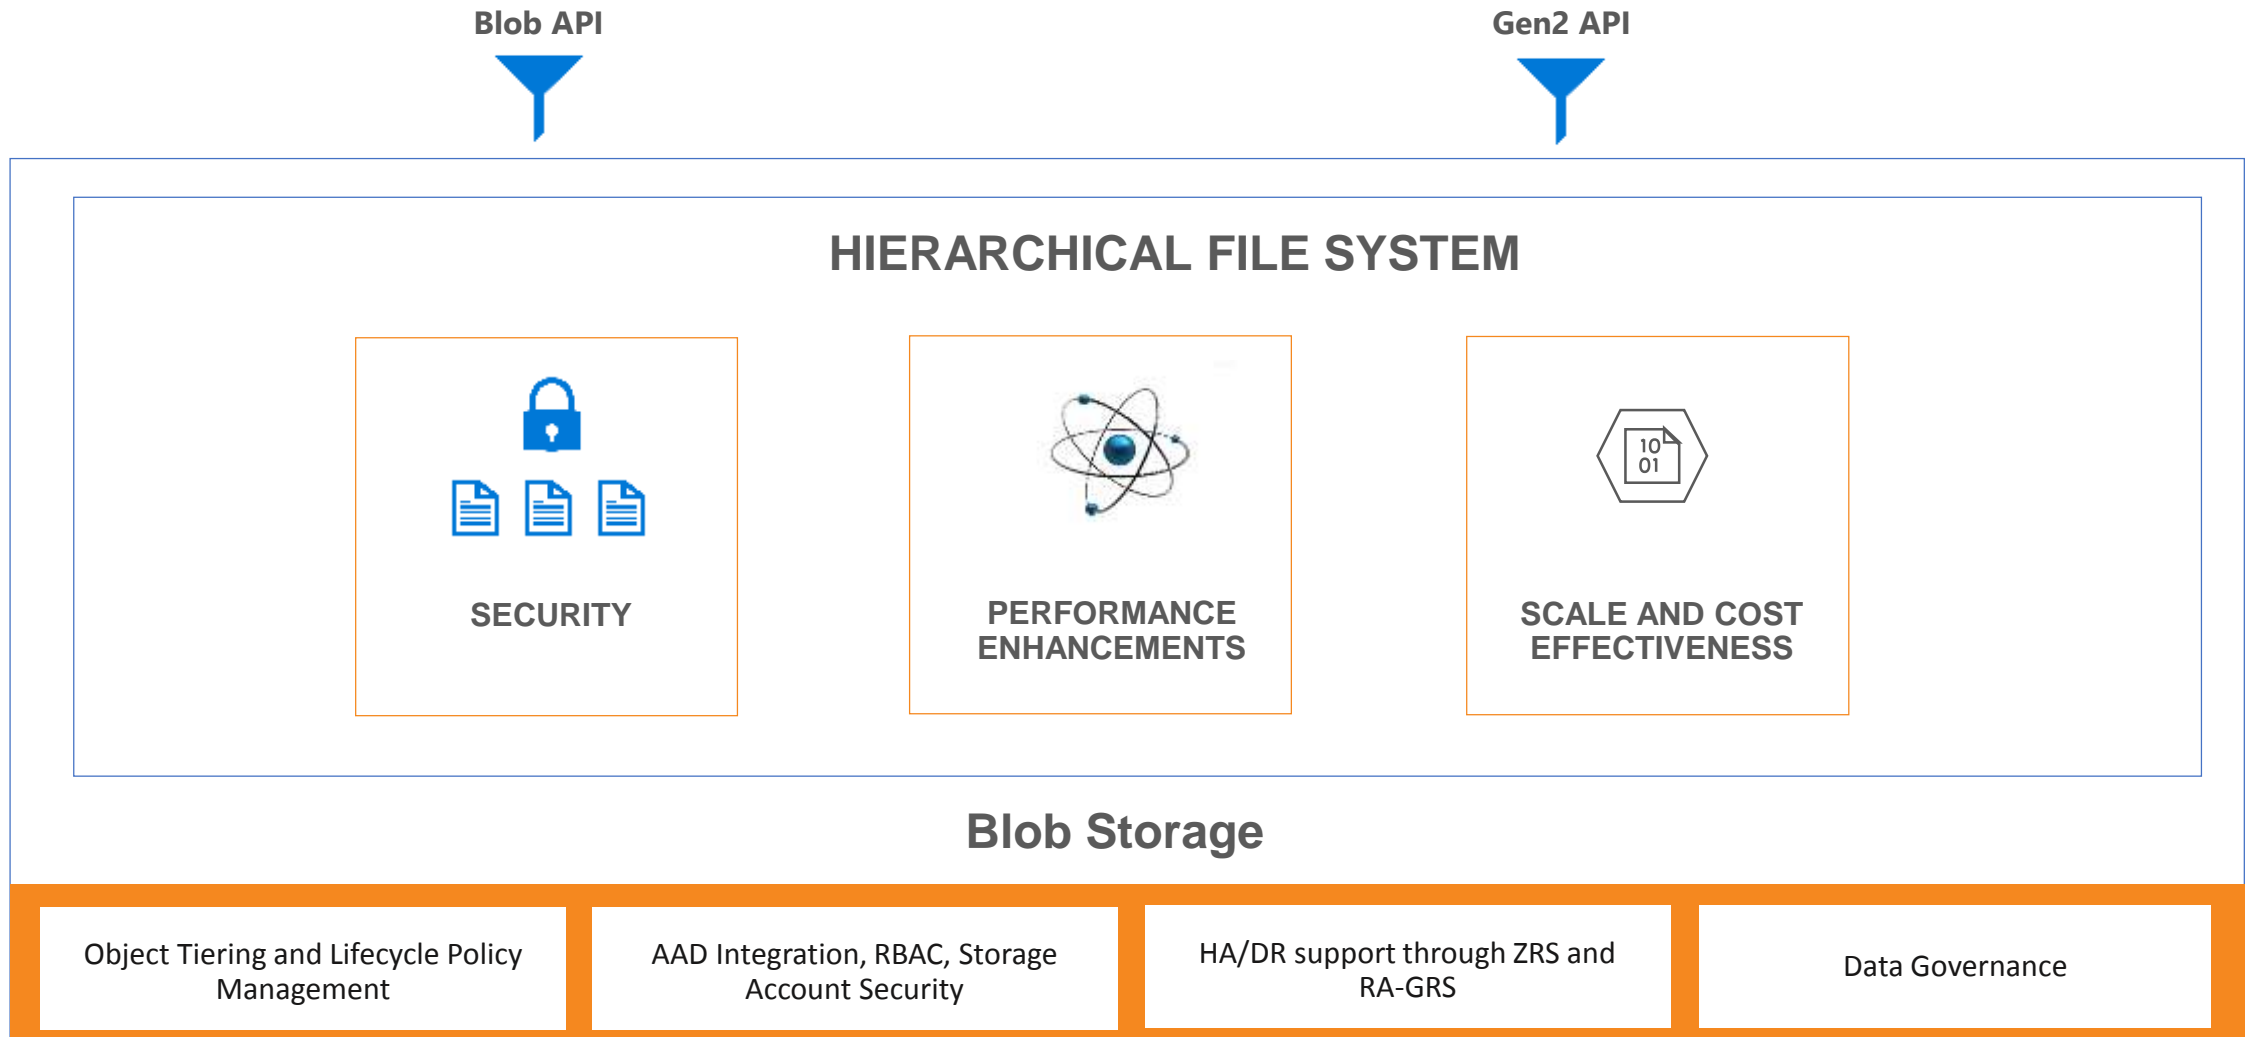
Object Store and File System access over the same data at the same time



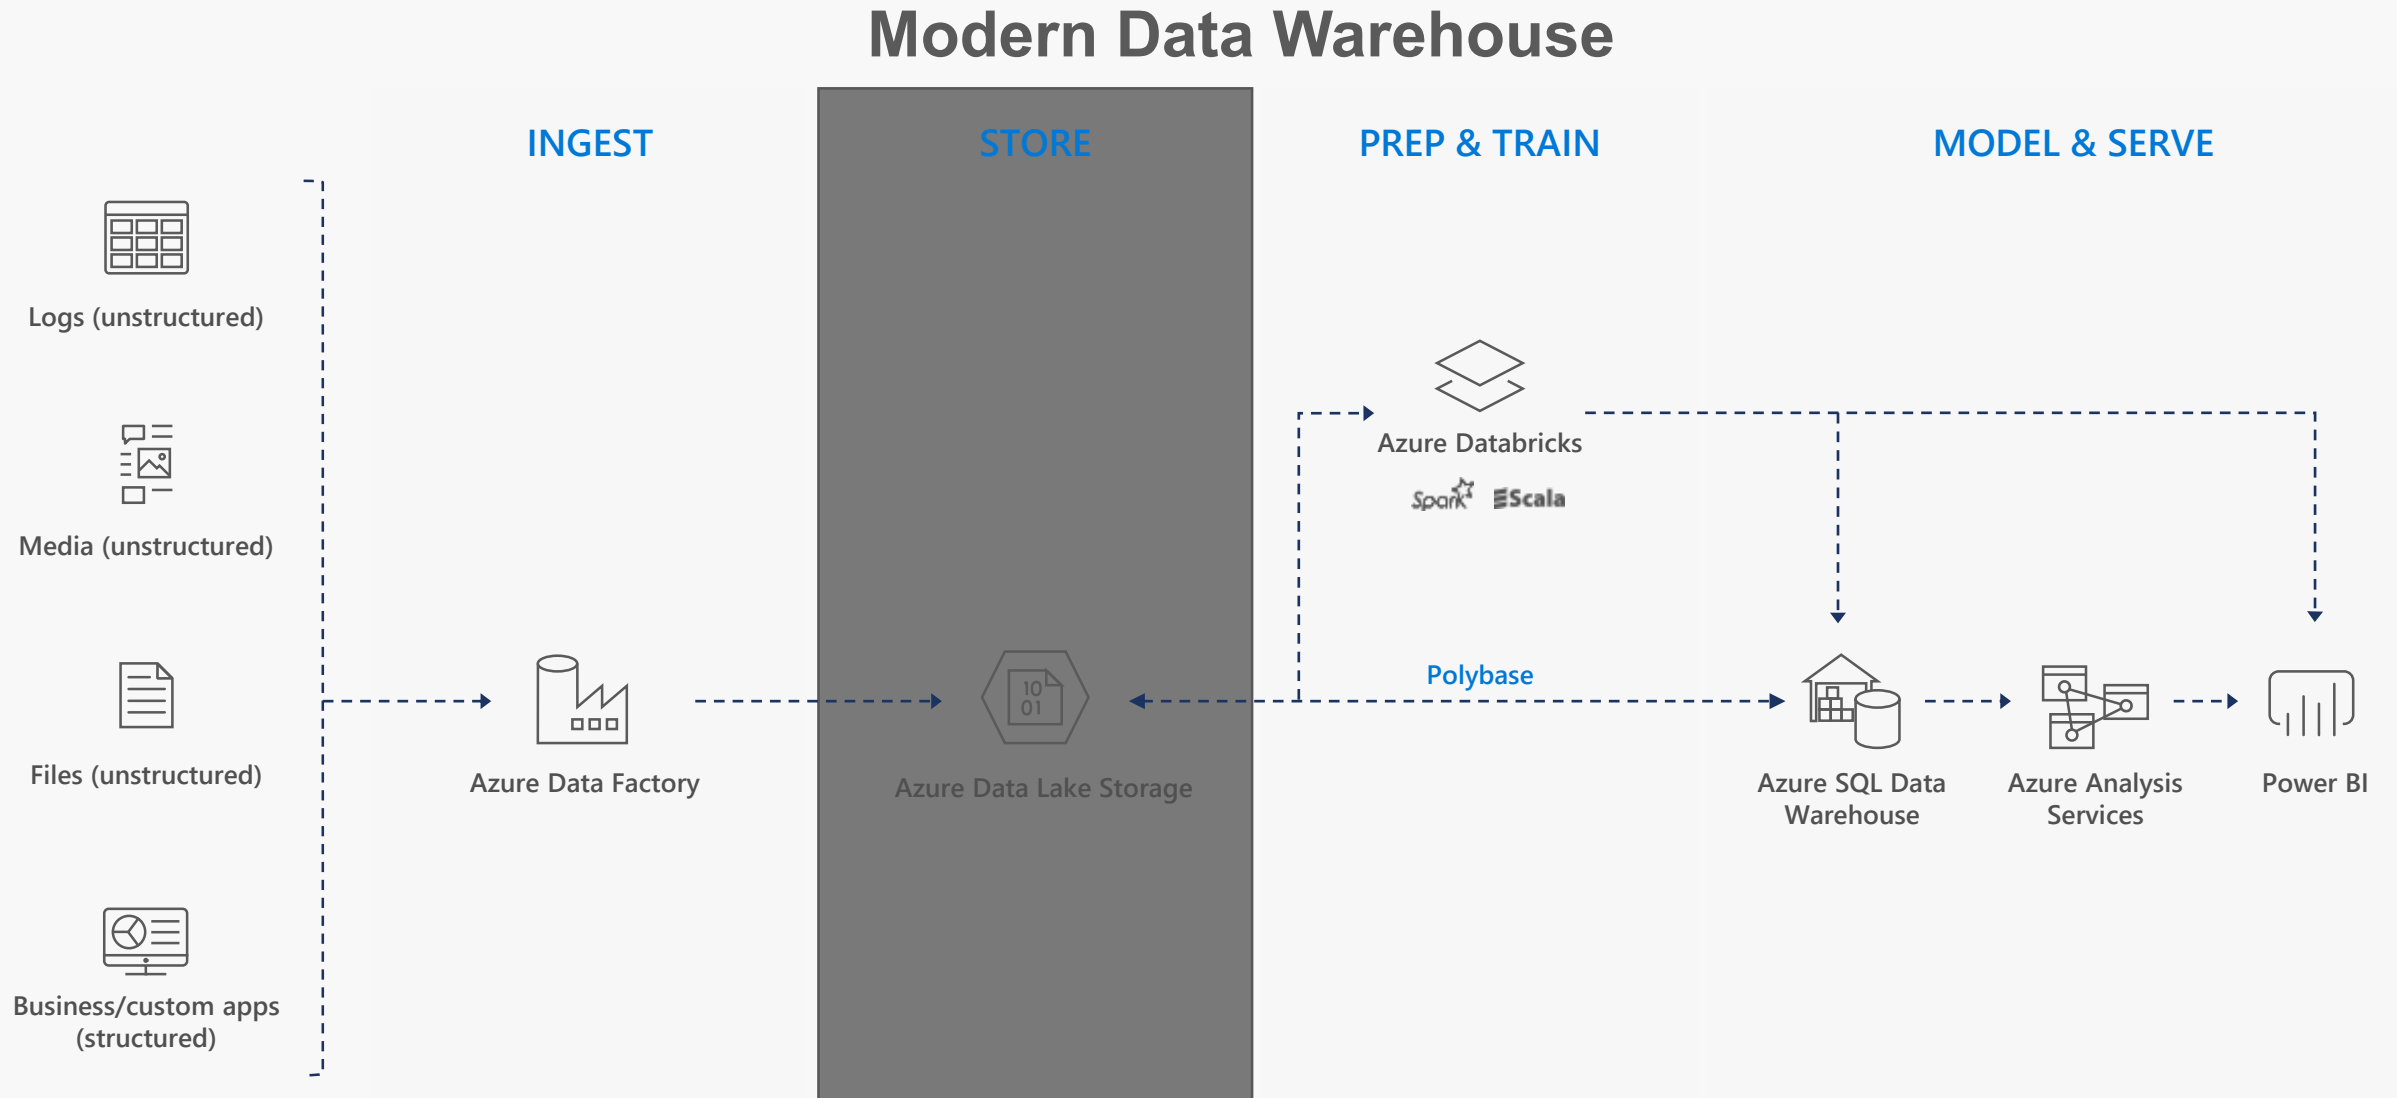
Azure Data Lake
Storage Gen2



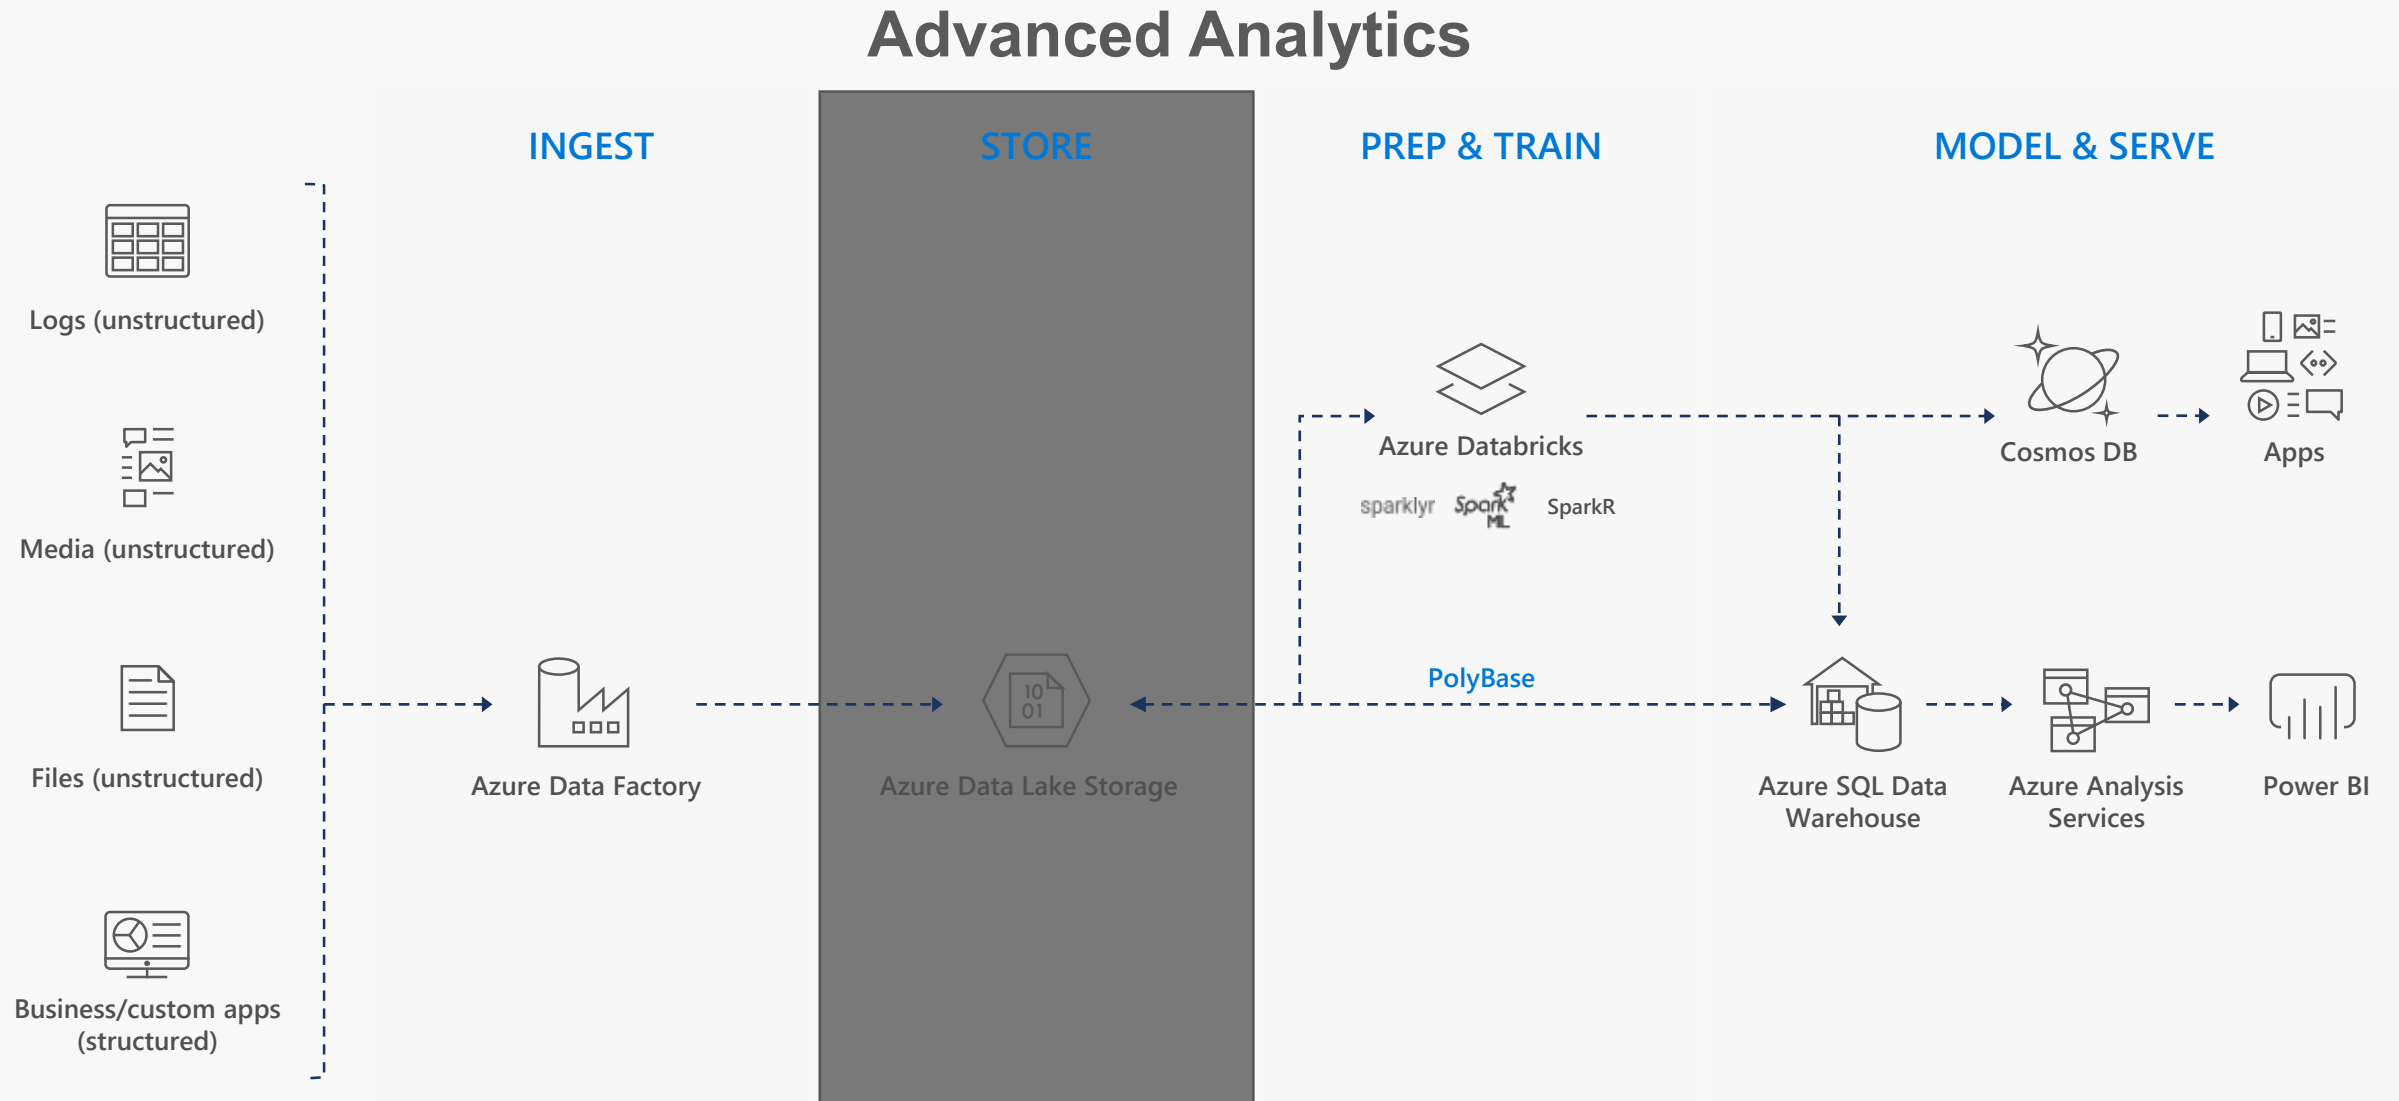
Azure Data Lake Storage Gen2 - Architecture



End-to-End Analytics

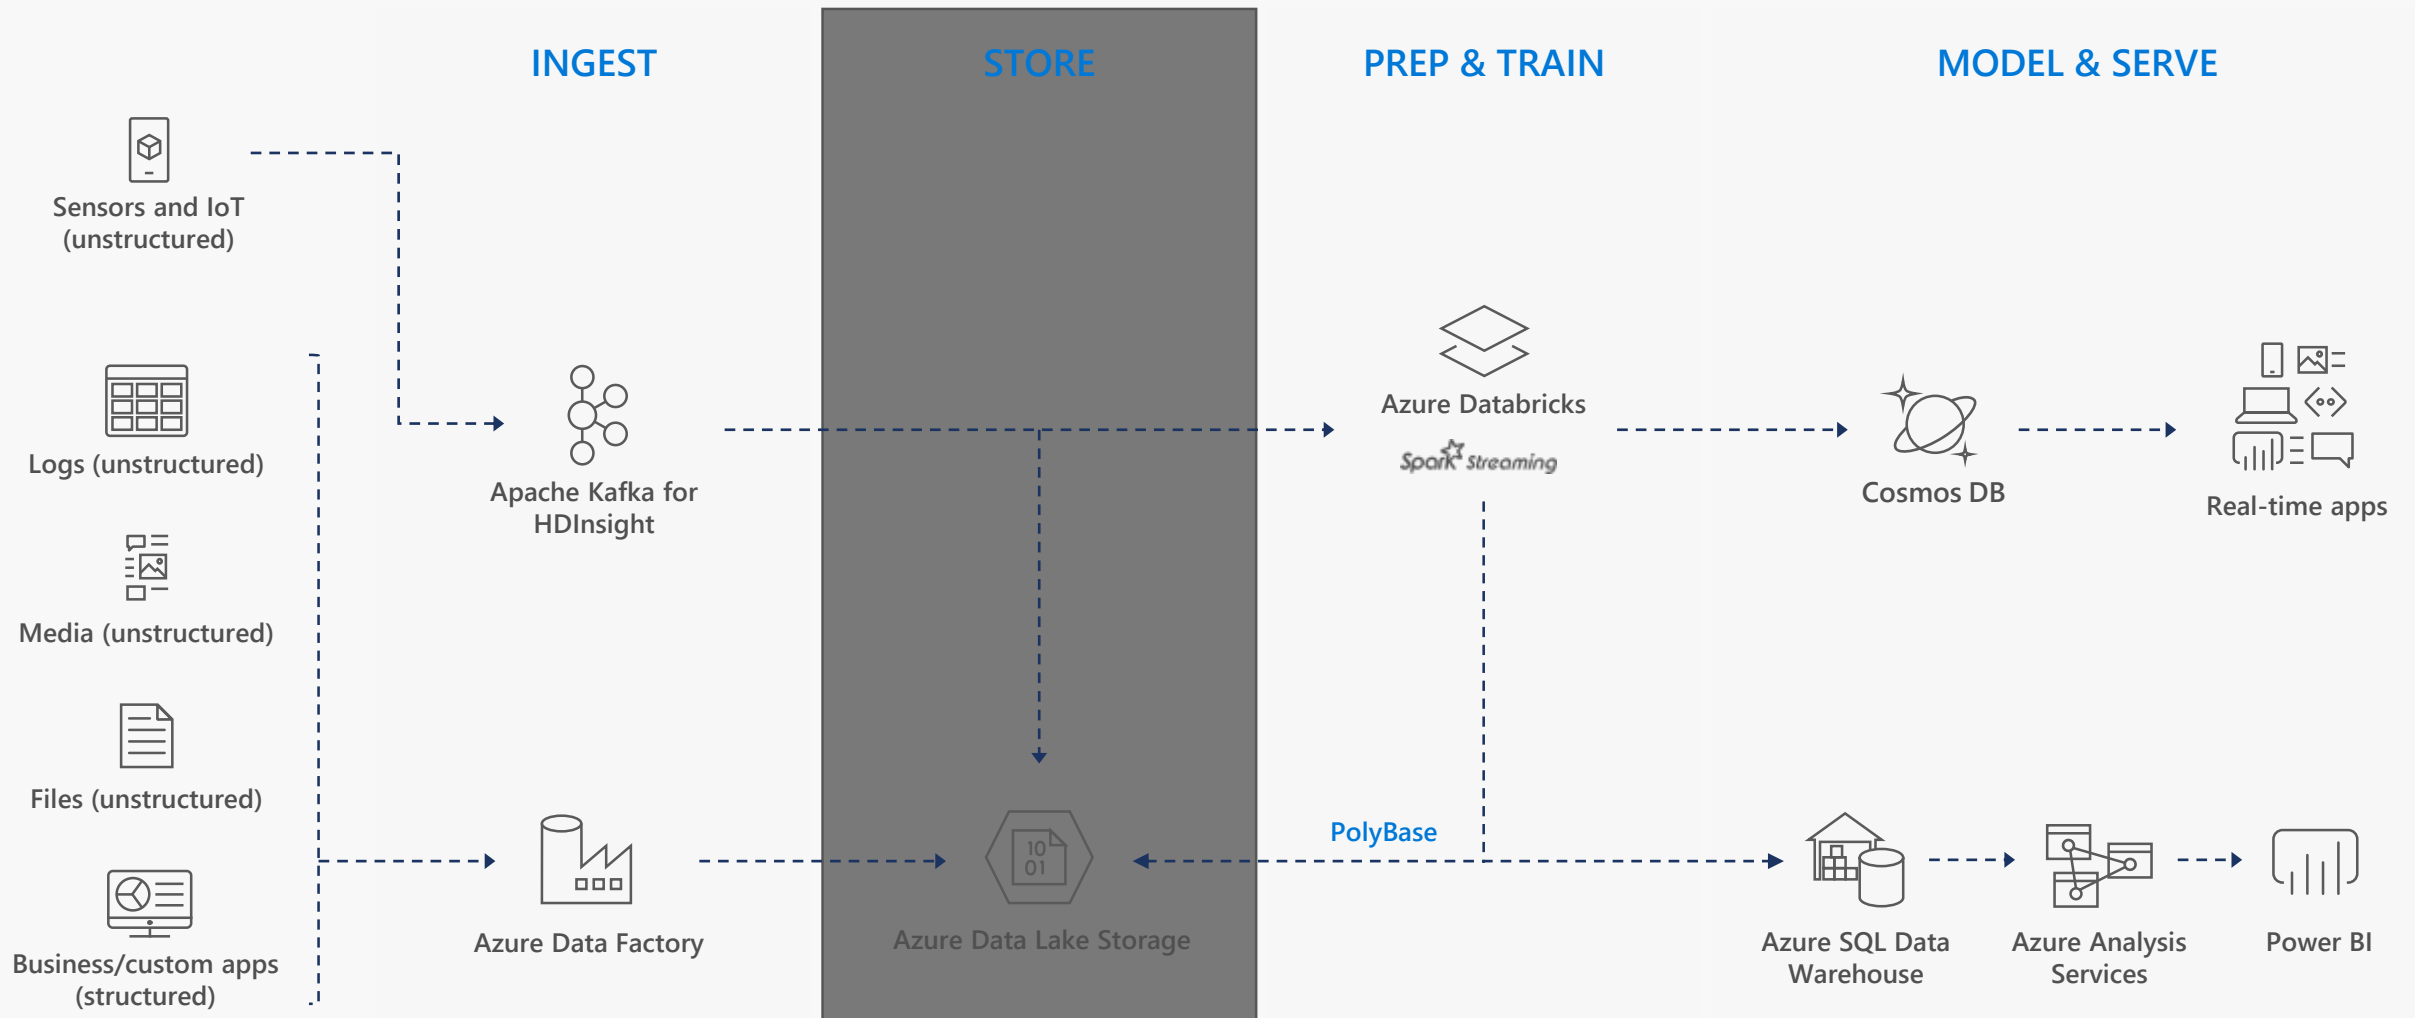


End-to-End Analytics



End-to-End Analytics

Realtime Analytics



Partner Ecosystem – Industry Verticals

- **Financial Services & Insurance**

- Compliance Archiving Insurance Risk Modeling
- Electronic Document Storage Financial Risk Modeling

- **Healthcare & Life Sciences**

- Genomics Medical Imaging (PACS, VNA)

- **Media & Entertainment**

- Media Asset Management Post Production
- Media Storage and Archiving Distribution

- **Public Safety & National Security**

- National Security (Body Cam) Enterprise Video Surveillance

- **Automotive**

- Autonomous Driving/Assisted Driving Connected Cars

- **Oil and Gas AEC Govt Retail**



Partner Ecosystem - Horizontals

Data Center Backup and Recovery



Distributed File Systems and Object Storage



Cloud Storage Gateway



Data Integration

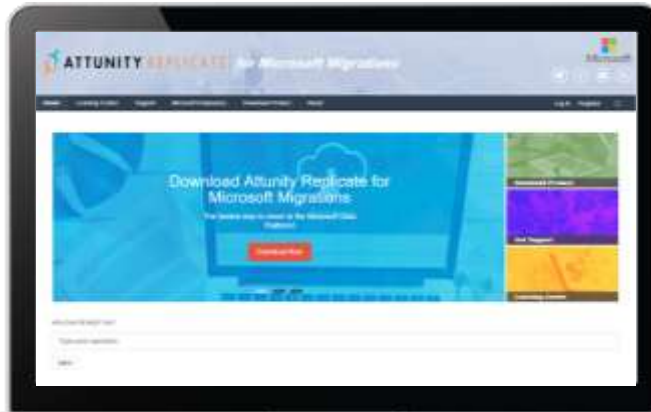


Attunity Replicate for Microsoft Migrations



FREE
DOWNLOAD

<https://aka.ms/attunity-replicate>



JOIN THE
COMMUNITY

<http://attunity.com/MicrosoftMigrations>



ASK
QUESTIONS

attunity-replicate@microsoft.com

Thank you

