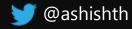
# Data Architects guide for successful Open Source analytics workloads with Azure HDInsight

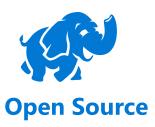
Ashish Thapliyal Principal Product Manager Azure HDInsight Microsoft Corp.





## **Azure HDInsight**

A secure and managed Apache Hadoop and Spark platform for building data lakes in the Cloud



- 100% Apache Open Source
- The most popular open source frameworks
- Part of the Hortonworks HDP distribution



#### Managed

- 99.9% availability SLA
- Cluster Health Monitoring
- Integration with Azure Log Analytics
- Highly optimized for Azure

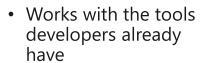


#### **Secure & Compliant**

- Role based access control
- Azure AD & Kerberos based authentication
- Strong VNET and service endpoint support
- The most trusted and compliant platform



#### **Productive**



 Special extensions for advanced debugging and diagnostics



#### **Lift & Shift**

- Move workloads from on-prem or other clouds without code changes
- Curated application platform for wide variety of use cases













# HDInsight Solution Architecture: Keys topics for today

OSS Framework Choices

Storage

HA & DR

Security

Monitoring

Cost Optimization

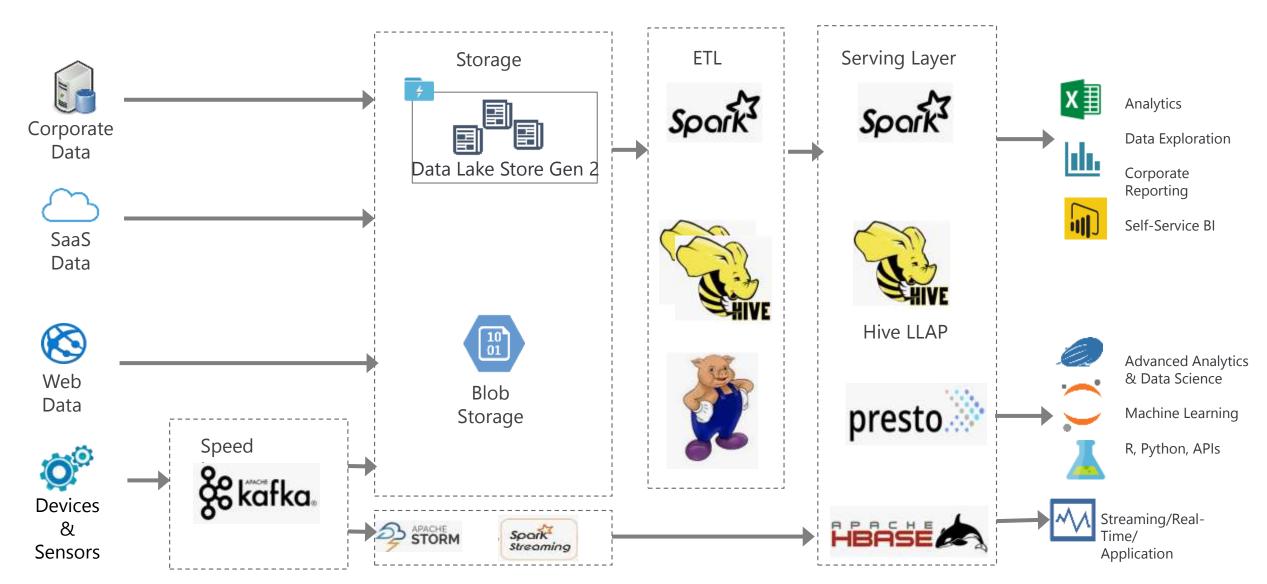


# 





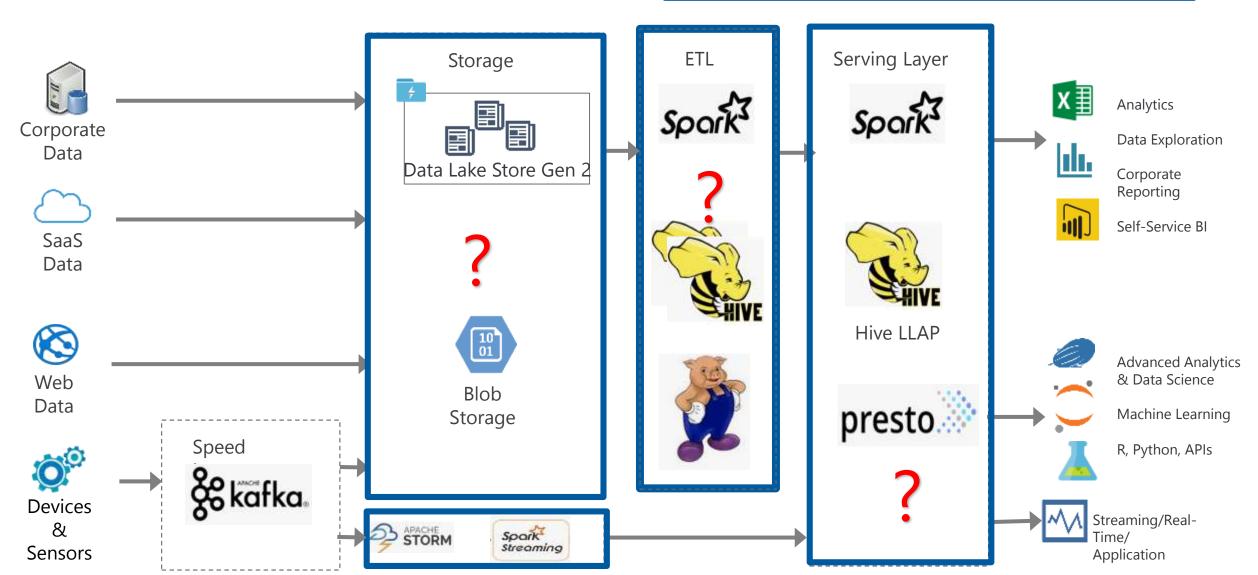






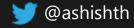
# 







# Scenario/ Technology mapping



# 1. ETL technology choices

	Spark	Pig	Hive
Designed for	ETL	ETL	Data warehousing
Adoption	High, increasing	Low, decreasing	Stable
Number of connectors	Highest	High	High
Languages	Python, R, Scala, Java, SQL	Pig	SQL
Performance	High	Medium	Medium



# 2. Streaming engine technology choices

	Spark Structured Streaming	Storm
Adoption	High, increasing	Decreasing
<b>Event processing guarantee</b>	Exactly once	At least once
Throughput	High	Low
<b>Processing Model</b>	Micro Batch	Real-Time
Latency	High	Low
<b>Event time support</b>	Yes	Yes
Languages	Python, R, Scala, Java, SQL	Java

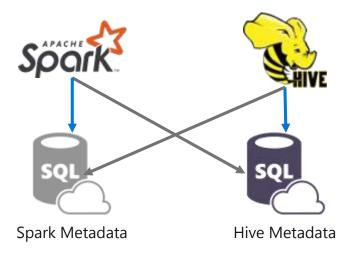


3.Interactive Query technology choices

s.interactive Query technology choices						
Capability	Hive LLAP	Spark SQL	Presto			
Interactive Query Speed	High	High	Medium			
Scale	High	High	Low			
Caching	Yes	Yes	Early Support			
Result Caching	Yes	No	No			
Intelligent Cache Eviction	Yes	No	No			
<b>Materialized Views</b>	Yes	No	No			
Complex Fact to Fact Joins	Yes	Yes	No			
Transactions	Yes	No	No			
Query Concurrency	High	Low	Low			
Row , Column level security	Yes [Apache Ranger+ AAD]	Medium	Medium			
Rich end user Tools	Yes	Yes	Yes			
Language Support	SQL, UDF	SQL, Scala, Python	SQL			
Data Source Connector Support	Storage Handlers	Data Sources	High number of connectors			

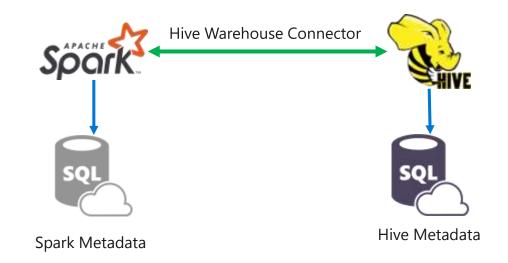
#### Tip: Spark & Hive Metastore

Azure HDInsight 3.6 with Hadoop 2.6



- Spark executors talk directly to Hive Metastore
- Reliability and compatibility issues
- Cannot take advantage of the native query engine

#### Azure HDInsight 4.0 with Hadoop 3.x



- New Hive Warehouse Connector
- Apache Arrow based communication between Spark executors and Hive LLAP
- Smart predicate pushdown
- Transactional access to Hive tables from Spark

**Hive Metastore migration tool**: https://azure.microsoft.com/en-us/blog/hdinsight-metastore-migration-tool-open-source-release-now-available/

# 4. Data pipeline orchestration technology choices

	ADF	Airflow	Oozie
Service management	Azure PaaS	laaS VM	HDInsight
Code	JSON	Python	Java
GUI	ADF V2 has great UX	Good UX	Below Average UX
Community	Microsoft	Growing (10893 Stars)	Declining (454 Stars)
On-demand clusters	On-demand clusters Yes		No
Extensibility	Custom action-only	Full, graph + actions	Custom action-only
Pipeline definition JSON/UX		Python/ UX	XML/UX
Devops-first design Yes		Yes	Yes
Pipeline monitoring Yes		Yes	Yes
Scheduling Event, Time		Event	Event, Time



## Migrating to Azure HDInsight Guide!

Motivation and benefits covers the benefits of migrating on-premises Hadoop ecosystem components to HDInsight and how to plan for the migration.

<u>Architecture best practices</u> provides best practices for the architecture of HDInsight systems and addresses different types of workloads.

<u>Infrastructure best practices</u> goes into detailed recommendations for managing the infrastructure of HDInsight clusters.

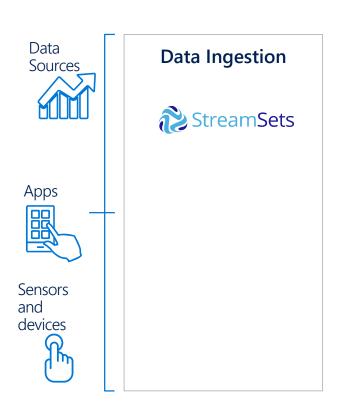
<u>Storage best practices</u> gives recommendations for data storage in HDInsight systems.

<u>Data migration best practices</u> provides recommendations for data migration to HDInsight.

<u>Security and DevOps best practices</u> gives recommendations for security and DevOps in HDInsight systems.

https://azure.microsoft.com/en-us/blog/migrating-on-premises-hadoop-infrastructure-to-azure-hdinsight/

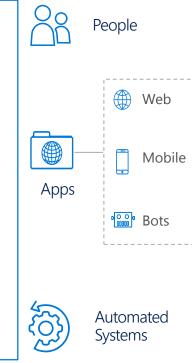
## HDInsight Applications: Something for everyone











Productivity Tools 📳 🥃 🔞









Data catalog/ Governance/ Lineage



Connectors: JDBC, ODBC A Simba

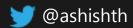








# Storage



#### **HDInsight Storage topics**

Data movement Storage options and tradeoffs

Caching



<b>Data Qty</b>	Network Bandwi	dth	
	45 Mbps (T3)	100 Mbps	1 Gbps
1 TB	2 days	1 day	2 hours
10 TB	22 days	10 days	1 day
35 TB	76 days	34 days	3 days
80 TB	173 days	78 days	8 days
100 TB	216 days	97 days	10 days
200 TB	1 year	194 days	19 days
500 TB	3 years	1 year	49 days
1 PB	6 years	3 years	97 days
2 PB	12 years	5 years	194 days



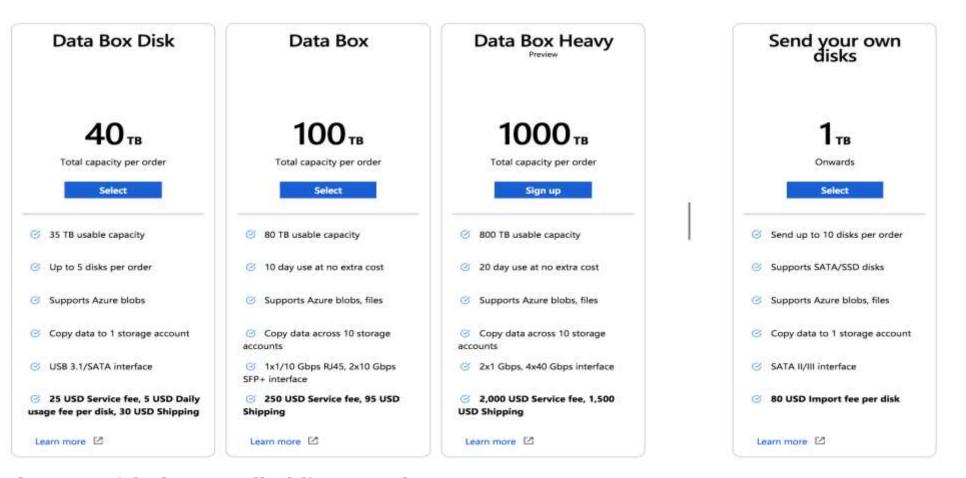
## **Storage Transfer options**

#### **Network Transfer with TLS**

- Over Internet
- Express Route
- Databox online Transfer
- Shipping data offline
- Import / Export service
- Data Box offline data transfer



#### **Data Box Transfer Options**



https://github.com/alkohli/azure-docs-pr/blob/4023eb52cc6ed103e0fa7e794e039c143b6d2a6a/articles/storage/blobs/data-lake-storage-migrate-on-prem-HDFS-cluster.md



# Storage options and tradeoffs



# **Storage Options with HDInsight**

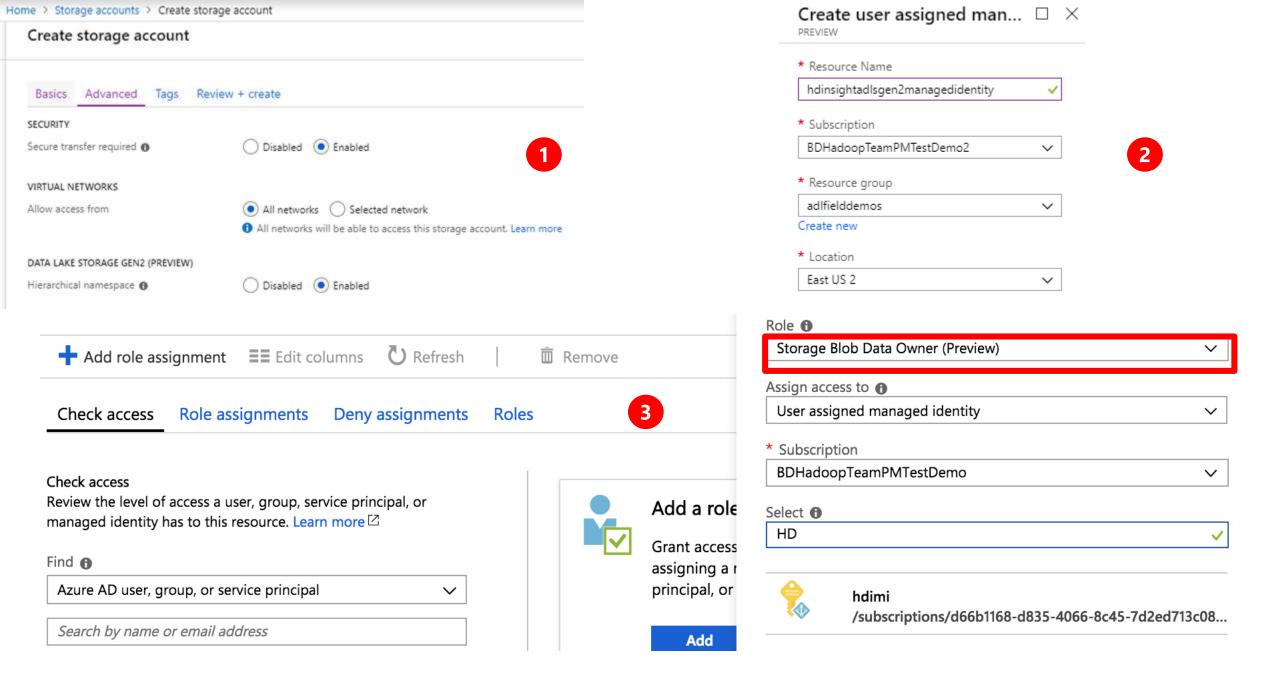
	Туре	Latency	Consistency	Workloads	Bandwidth	Key Benefits
ADLS Gen 1	Hierarchical	10-100ms	Low	HDInsight 3.6( No HBase)	High	Atomic Rename, File Folder level ACL's
ADLS Gen 2	Hierarchical	10-50ms	Medium	HDInsight 3.6 & 4.0	Unconstrained	Atomic Rename, File Folder level ACL's
Standard BLOB	Object Store	10-50ms	Medium	HDInsight 3.6 & 4.0	Unconstrained	Mature
Premium BLOB	Object Store	~5ms	High	HBase in Preview	Unconstrained	Fast
Premium Managed Disks	Hierarchical	~5ms	High	Kafka, HBase in preview	Based on disk	Consistent



## ADLS Gen 1 and Gen 2 & HDInsight Compatibility

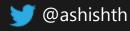
Scenario	Supported	Workaround
HDInsight 3.6 & 4.0 with <b>Standard Blob as Primary</b> and/ or secondary	Yes	
HDInsight 3.6 & 4.0 with ADLS Gen2 as primary	Yes	
HDInsight 3.6 & 4.0 with <b>ADLS Gen2 as primary &amp; Blob as additional</b>	Yes	
HDInsight 3.6 & 4.0 with <b>Blob as primary &amp; ADLS Gen2 as additional</b>	No	
HDInsight 3.6 with multiple <b>ADLS Gen2 accounts</b>	Yes	
HDInsight 3.6 & 4.0 with ADLS Gen1 and ADLS Gen 2	No	Distcp across two clusters
HDInsight 4.0 with ADLS Gen 1	No	Distcp across two clusters



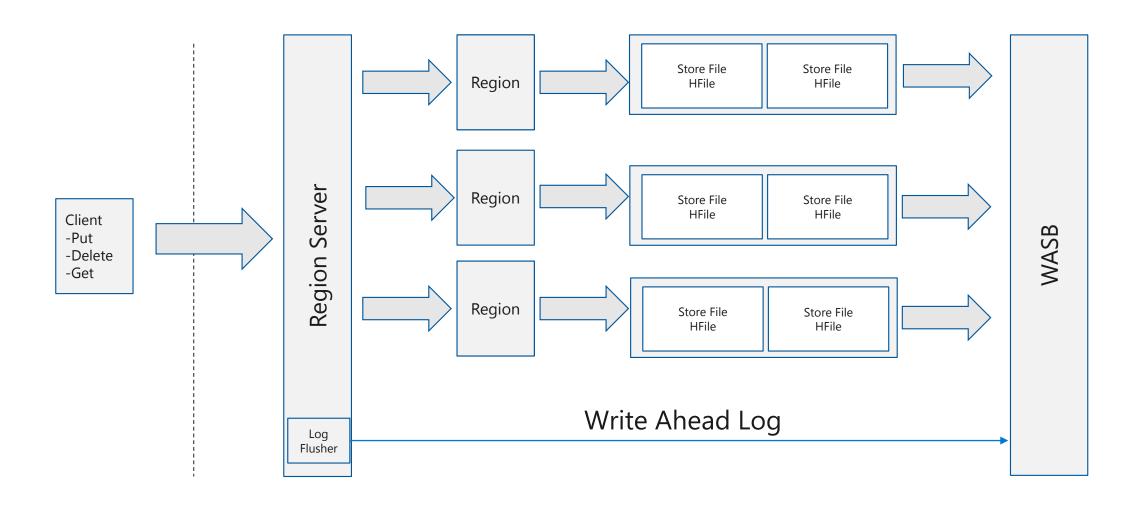


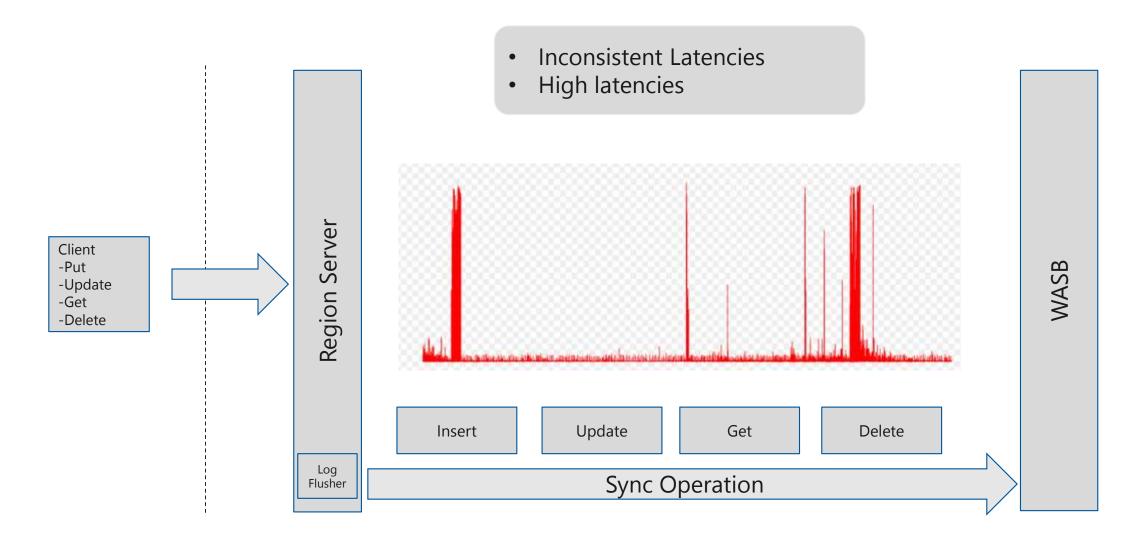


New Storage Improvements: Enhanced reads and writes for low latency workloads (Private Preview)



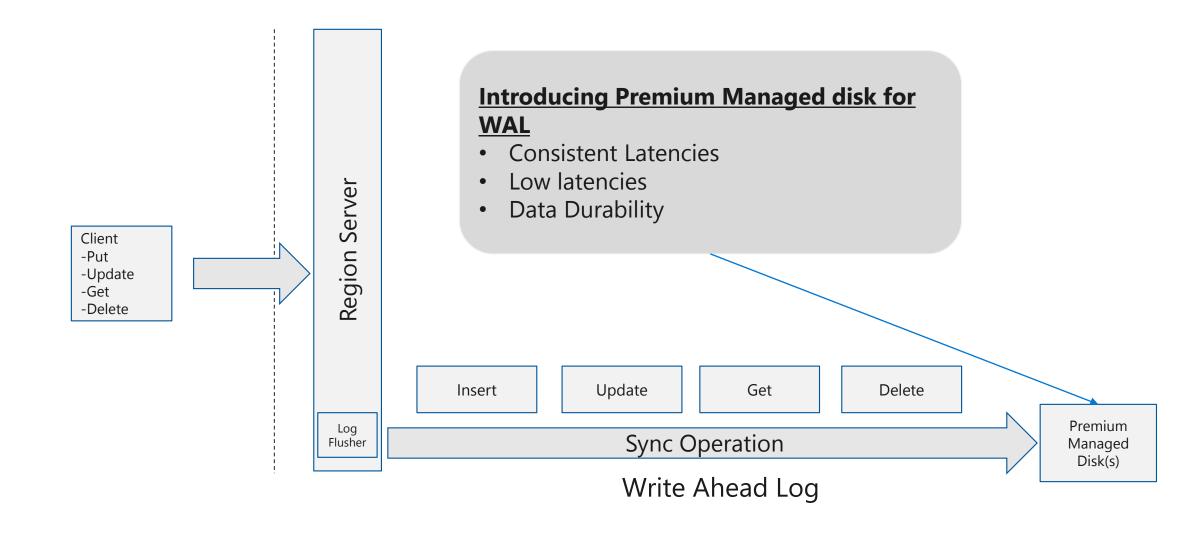
## Example: Low latency workload HBase/ Small write



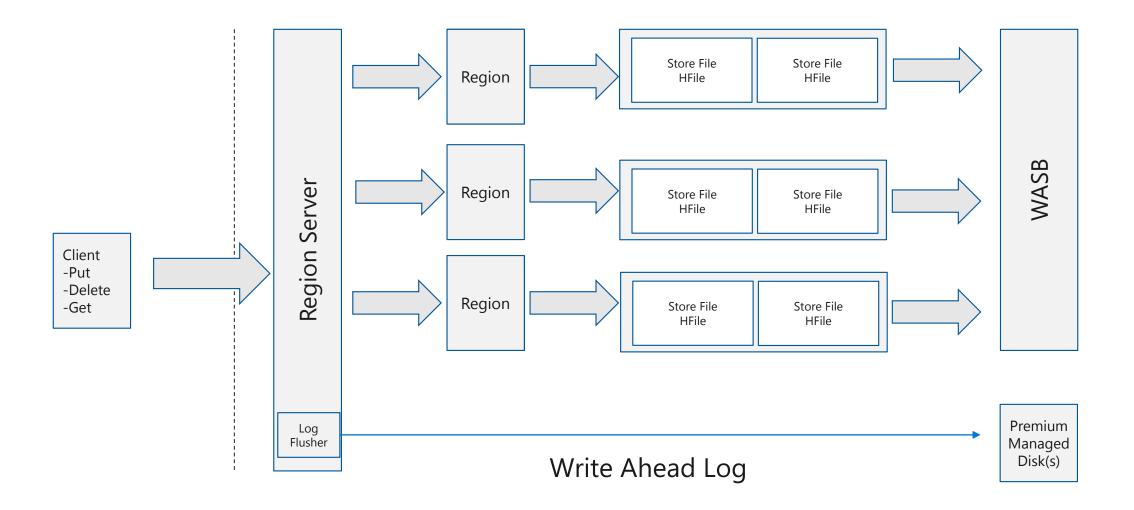


Write path challenges with Write Ahead Log



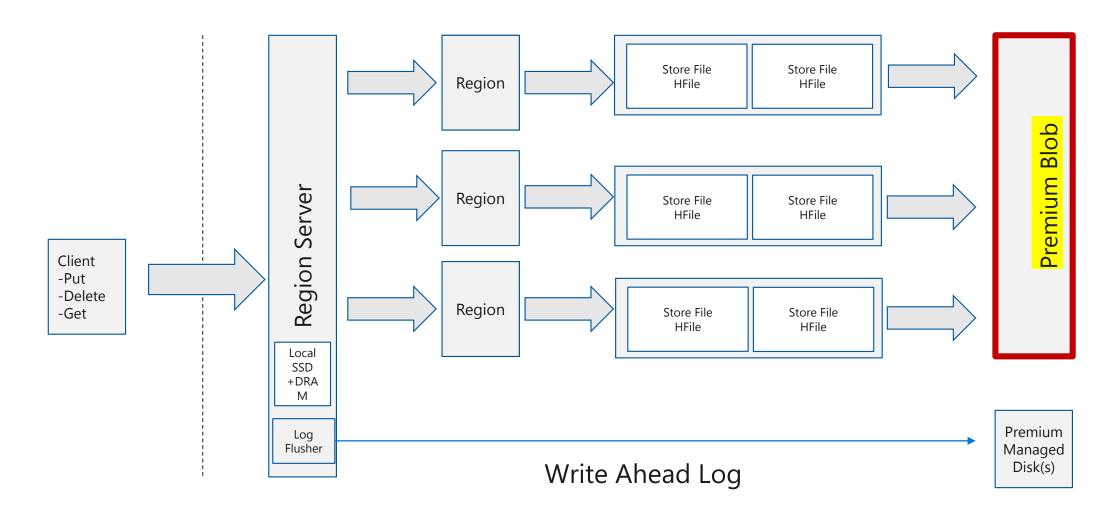






#### **How about Reads?**





# **Introducing support for Premium Blob**



Performance (YCSB)

Cluster Type	Operation	Row Size	# ops	#Region Servers	Region Server Node Size	#Clients	Throughput	Avg Latency (ms)	Run Time (min)
Standard	Write	1KB	107,374,182	4	Standard_D4_V2	2	37,958	0.417	47
Premium WAL	Write	1KB	107,374,182	4	Standard_D4_V2	2	57,812	0271	31
Standard	Small Write	100 Bytes	1,073,741,824	4	Standard_DS4_V2	2	84,910	0186	210
Premium WAL	Small Write	100 Bytes	1,073,741,824	4	Standard_DS4_V2	2	701,234	0.016	25
Standard	Read	100 Bytes	925,075	4	Standard_D4_V2	2	256	62	60
Premium WAL & Premium Blob	Read	100 Bytes	33,503,676	4	Standard_D4_V2	2	9,306	1.7	60
Standard	Large Read	1K	945,682	4	Standard_D4_V2		262	61	60
Premium WAL & Premium Blob	Large Read	1K	24,846,209	4	Standard_D4_V2	2	6901	2.3	60

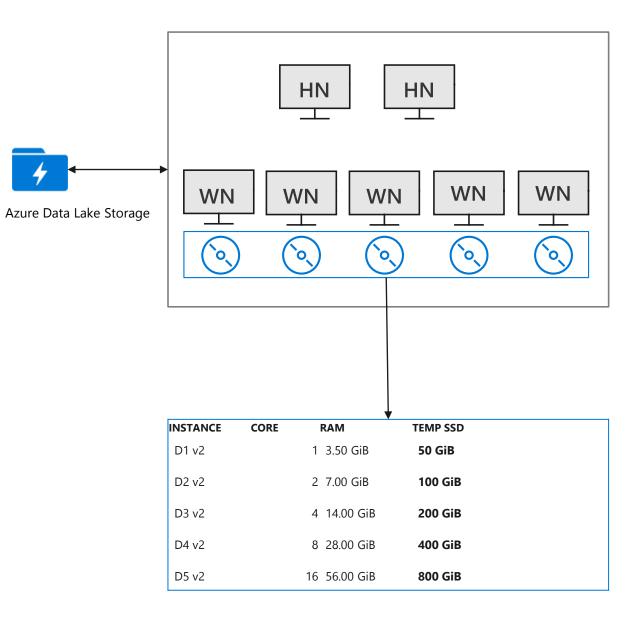
# Caching options in HDInsight



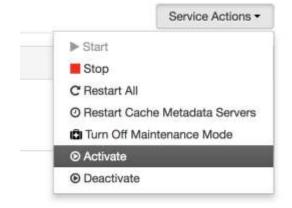
# **Caching Options by workload**

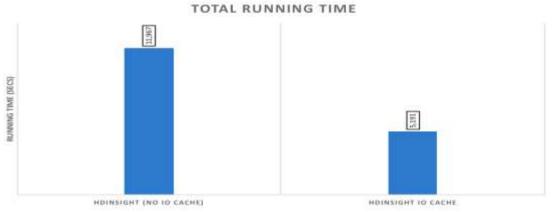
Workload	Caching Options	Key benefits
Spark	Spark IO Cache	Up to ~8 to 10x perf improvements
HBase & Phoenix	Bucket cache	Up 5-10x perf gains on recently read or written data
Hive + LLAP	LLAP Intelligent cache/Result Cache	Up to ~4-100X gain on cached data

## New HDInsight IO Cache

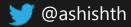


- Significant Spark performance speed up with IO cache (up to 9X perf gains)
- Automatic cache resource management
- DRAM + Temp SSD makes large cache

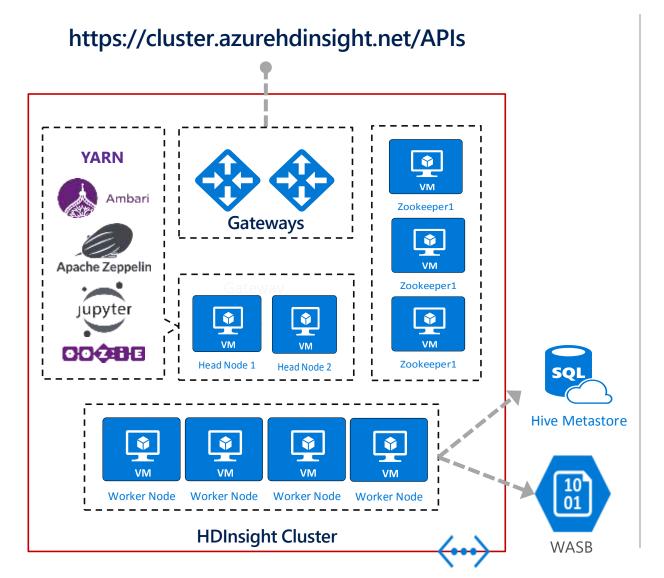




## HA & DR



## **Azure HDInsight: Highly Available End-points**



#### **Highly Available APIs:**

Livy – Spark job submission and interactive session management

Yarn – cluster resource management, Yarn job submission

Ambari – cluster management

Oozie – Oozie workflow scheduling and coordination (legacy APIs, ADF is recommended as a replacement)

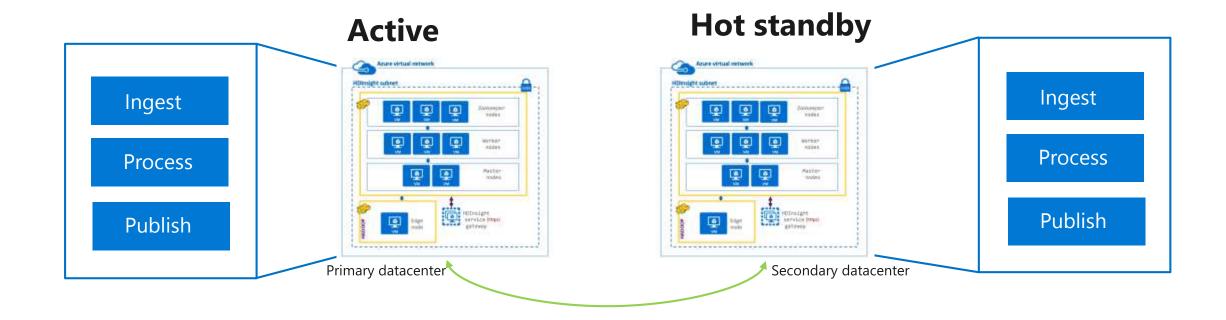


# DR options by workloads

Workload	DR Option
Spark / Hive	Manual, Partner solution
HBase	HBase replication, Snapshot export, Import Export, Copy Tables
Kafka	Mirror Maker



## Option 1: Active-hot standby with dual ingest



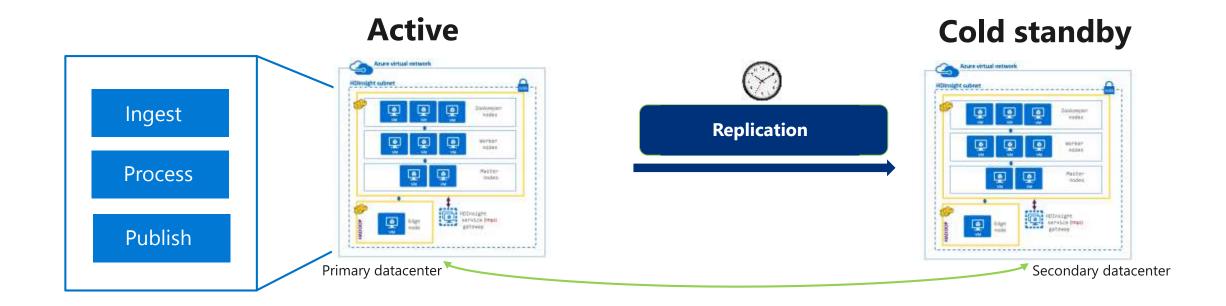
RPO RTO Cost Low None High

- Applications/integration processes write to both clusters
- Both clusters run identical batch jobs
- Standby cluster is <u>offline</u> for reads by applications and end users
- Synchronization tasks need to be run to ensure clusters are in sync

Its also not uncommon to have the hot standby used for non-critical workloads like adhoc exploratory analytics by data analysts, for reporting



## Option 2: Active- cold standby with batch replication

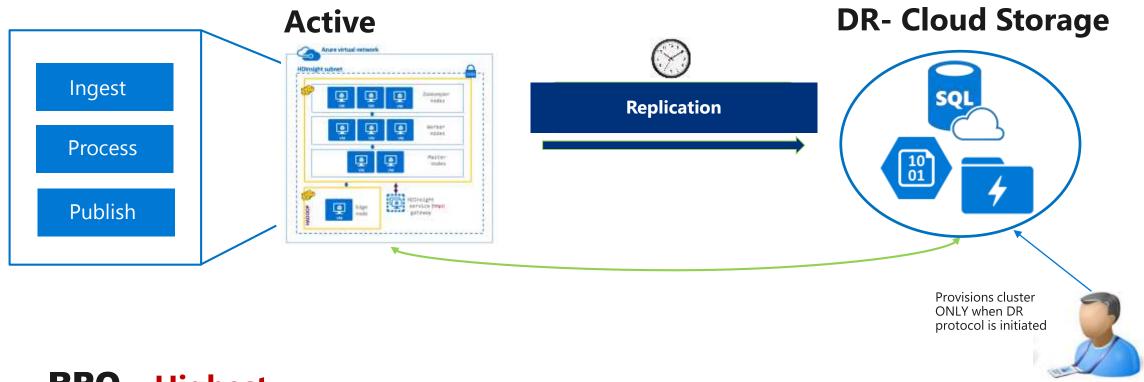


RPO Medium RTO Medium Cost High

- Applications write to active-primary ONLY
- Replication to DR cluster is incremental, batch, scheduled
- Synchronization tasks need to be run to ensure clusters are in sync
- Its is not uncommon to have a storage dense, compute light DR cluster for cost optimization
- Only curated data is replicated
- No batch processes are run



## Option 3: Active-DR on demand

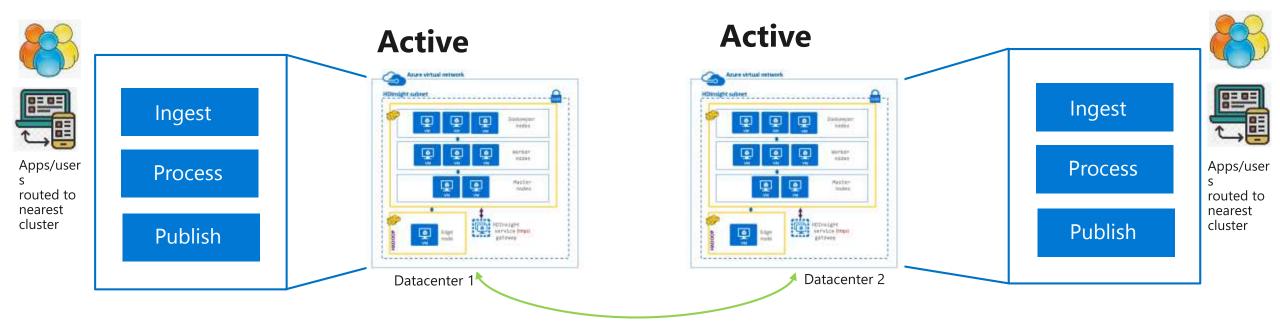


RPO Highest RTO Highest Cost Lowest Effort Highest

- Applications write to active-primary cluster, no DR cluster provisioned
- Replication utilities synchronize data and metastore to cloud storage in DR datacenter
- Administrator provisions cluster on-demand in DR datacenter
- Synchronization tasks need to be run to ensure clusters are in sync



## Option 4: Active-Active with multi-master (WANdisco)



RPO RTO Cost Lowest None Highest

- Applications/integration processes write to/read from nearest cluster clusters
- Both clusters run identical batch jobs
- Synchronization tasks need to be run to ensure clusters are in sync
- WANdisco fusion supports automated metadata replication (sync) and data replication (async), scheduling, monitoring, alerting, bandwidth throttling and more

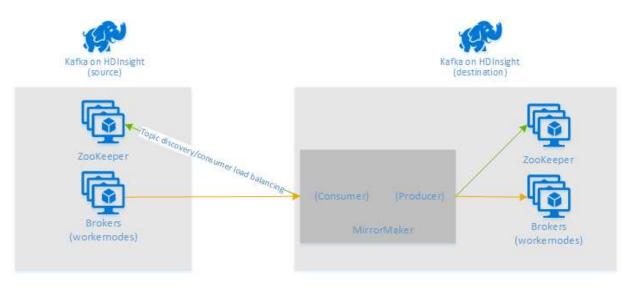


## Kafka Replication with MirrorMaker



- Kafka is often deployed in multiple environments for Disaster Recovery, high availability, and onprem to cloud hybrid scenarios.
- These require replication of data from one Kafka to the other. HDInsight has worked closely with enterprise customers to understand this need, and provides support for data replication scenarios through Apache MirrorMaker





## More Resources

Spark/ Hive HA & DR <a href="https://github.com/anagha-microsoft/hdi-spark-dr">https://github.com/anagha-microsoft/hdi-spark-dr</a>

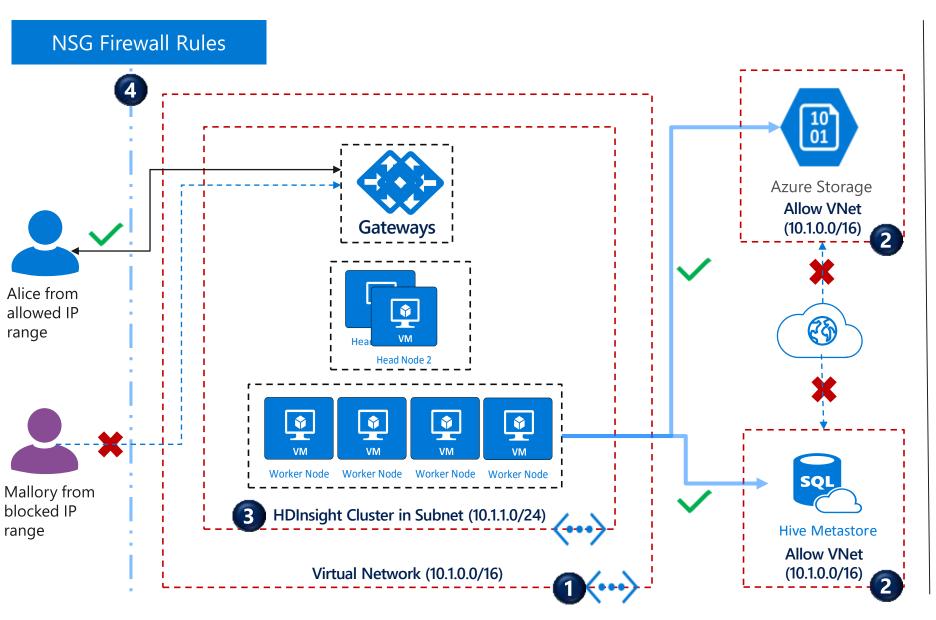
Kafka HA & DR <a href="https://github.com/anagha-microsoft/hdi-kafka-dr">https://github.com/anagha-microsoft/hdi-kafka-dr</a> HBase Backup, Replication <a href="https://docs.microsoft.com/en-us/azure/hdinsight/hbase/apache-hbase-backup-replication">https://docs.microsoft.com/en-us/azure/hdinsight/hbase/apache-hbase-backup-replication</a>

# Security



## **Azure HDInsight Network Security**

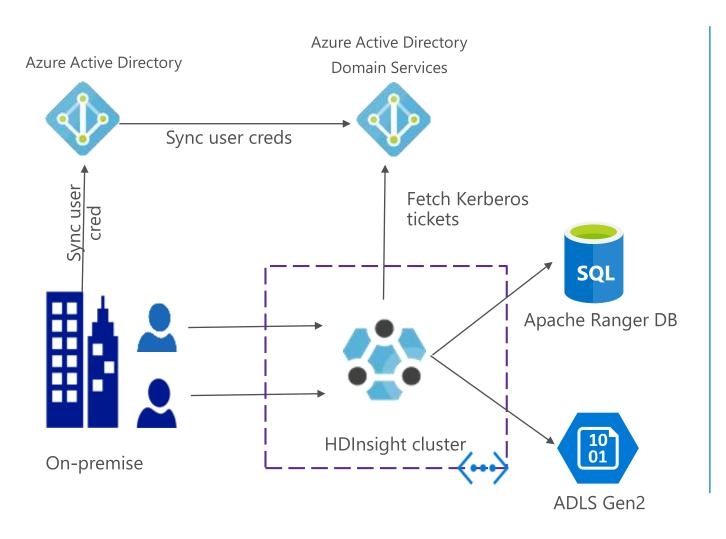
**Securing Data sources with Virtual Network Service Endpoints** 



- Create VNet, a subnet and enable service endpoint
- Restrict network access to Storage & SQL
- 3 Create HDInsight cluster within subnet
- Create NSG rules to control inbound access to HDInsight cluster



## **Azure HDInsight: Authentication & Access Control**



#### **Authentication:**

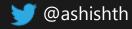
- Supports identities managed in Azure Active Directory (AAD)
- Clusters are joined to Active Directory
   Domain Services (ADDS) based Kerberos
   Domain Controllers.
- On-premise corporate identities are synced to AAD and ADDS via AD Federation Services.

#### **Access Control:**

- Apache Ranger based access control and auditing
- Ranger plugins for Hive, Spark, Kafka and HBase.

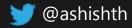


# Ranger + ADLS Gen 2 Auth Scenarios in HDInsight



Scenario	Authorizing Component
Yarn: Submit-App	Apache Ranger: Yarn Plugin
Hive Operations: Create, Select, Update, Drop, index, Lock, Read, Write, Masking, Row level filter on Hive Database, Table & Columns	Apache Ranger: Hive Plugin
Create/ Alter Table with storage location reference	Apache Ranger + ADLS Gen 2 ACL's
Spark SQL access with Hive Metastore	Apache Ranger: Hive Plugin
HBase Access Policies	Apache Ranger/ HBase plugin
Kafka Access Policies	Apache ranger/ Kafka Plugin
Access Azure Data Lake Storage Gen2 using the Spark DataFrame API	ADLS Gen 2 ACLs
Access Azure Data Lake Storage Gen2 using the RDD API	ADLS Gen 2 ACLs
HDFS operations: Mkdir, ls, put, copyFromLocal, get, cat, mv, cp etc	ADLS Gen 2 ACLs
Running Map Reduce jobs	ADLS Gen 2 ACLs

# Monitoring



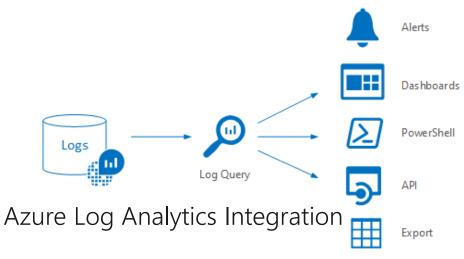
## **HDInsight Monitoring options**



- View cluster metrics like CPU, memory, and disk usage at a glance in real time
- Identify malfunctioning components with Ambari alerts
- Monitor queue capacities, jobs, and view associated OSS logs



- See gateway requests to monitor cluster stress and cluster size to monitor costs
- Apply filters and chart splitting to extract important data
- Set up alert rules to receive notifications and trigger actions for key metrics



- Organizes cluster metrics and OSS log records into queryable tables
- Create custom dashboards to surface all the metrics you need from multiple clusters on a single pane of glass

# **Cost Optimization**



## **Lower Price**

Different VM Pricing Options: HDInsight pricing model is VM price + HDInsight markup

### Different VM SKU Options

INSTANCE	VCPU		TEMPORARY STORAGE	PAYG
D14 v2	16	112.00 GiB	800 GiB	\$1.196/hour
E16 v3	16	128.00 GiB	400 GiB	\$1.064/hour

11%

## Different Region Options

• A8 V2 West US 2 region, HDI price: \$0.469/hour

12.5%

A8 V2 East US 2 region, HDI price: \$0.536/hour



## Thank You!

